



Serverless Messaging with Pub/Sub

Agenda


Processing Streaming Data

Cloud Pub/Sub

Cloud Dataflow Streaming
Features

BigQuery and Bigtable Streaming
Features

Advanced BigQuery
Functionality



Cloud Pub/Sub



Cloud
Pub/Sub

Qualities that Cloud Pub/Sub
contribute to Data Engineering
solutions:

Availability
Durability
Scalability

100s of
milliseconds

Cloud Pub/Sub



Cloud
Pub/Sub

Qualities that Cloud Pub/Sub
contribute to Data Engineering
solutions:

Availability
Durability
Scalability

100s of
milliseconds

Cloud Pub/Sub



Cloud
Pub/Sub

Qualities that Cloud Pub/Sub
contribute to Data Engineering
solutions:

Availability
Durability
Scalability

100s of
milliseconds

Cloud Pub/Sub




Cloud
Pub/Sub

Qualities that Cloud Pub/Sub
contribute to Data Engineering
solutions:

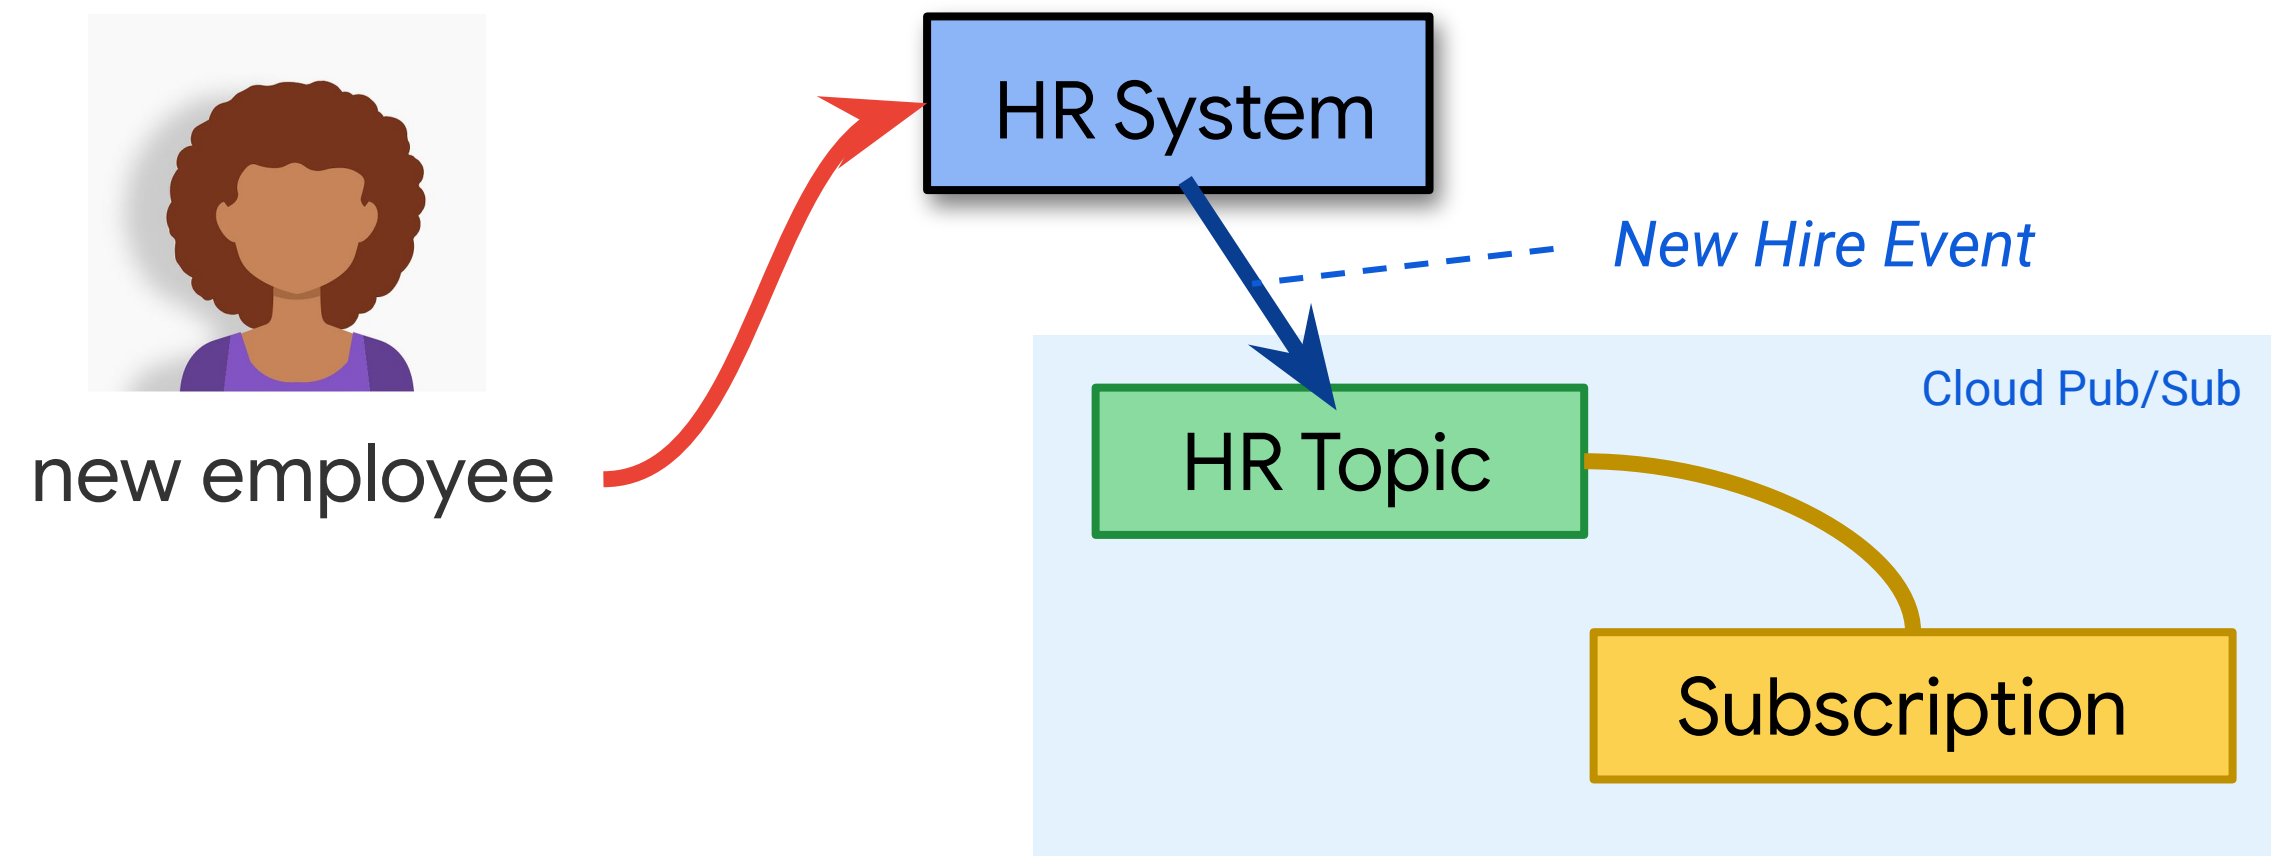
Availability
Durability
Scalability

100s of
milliseconds

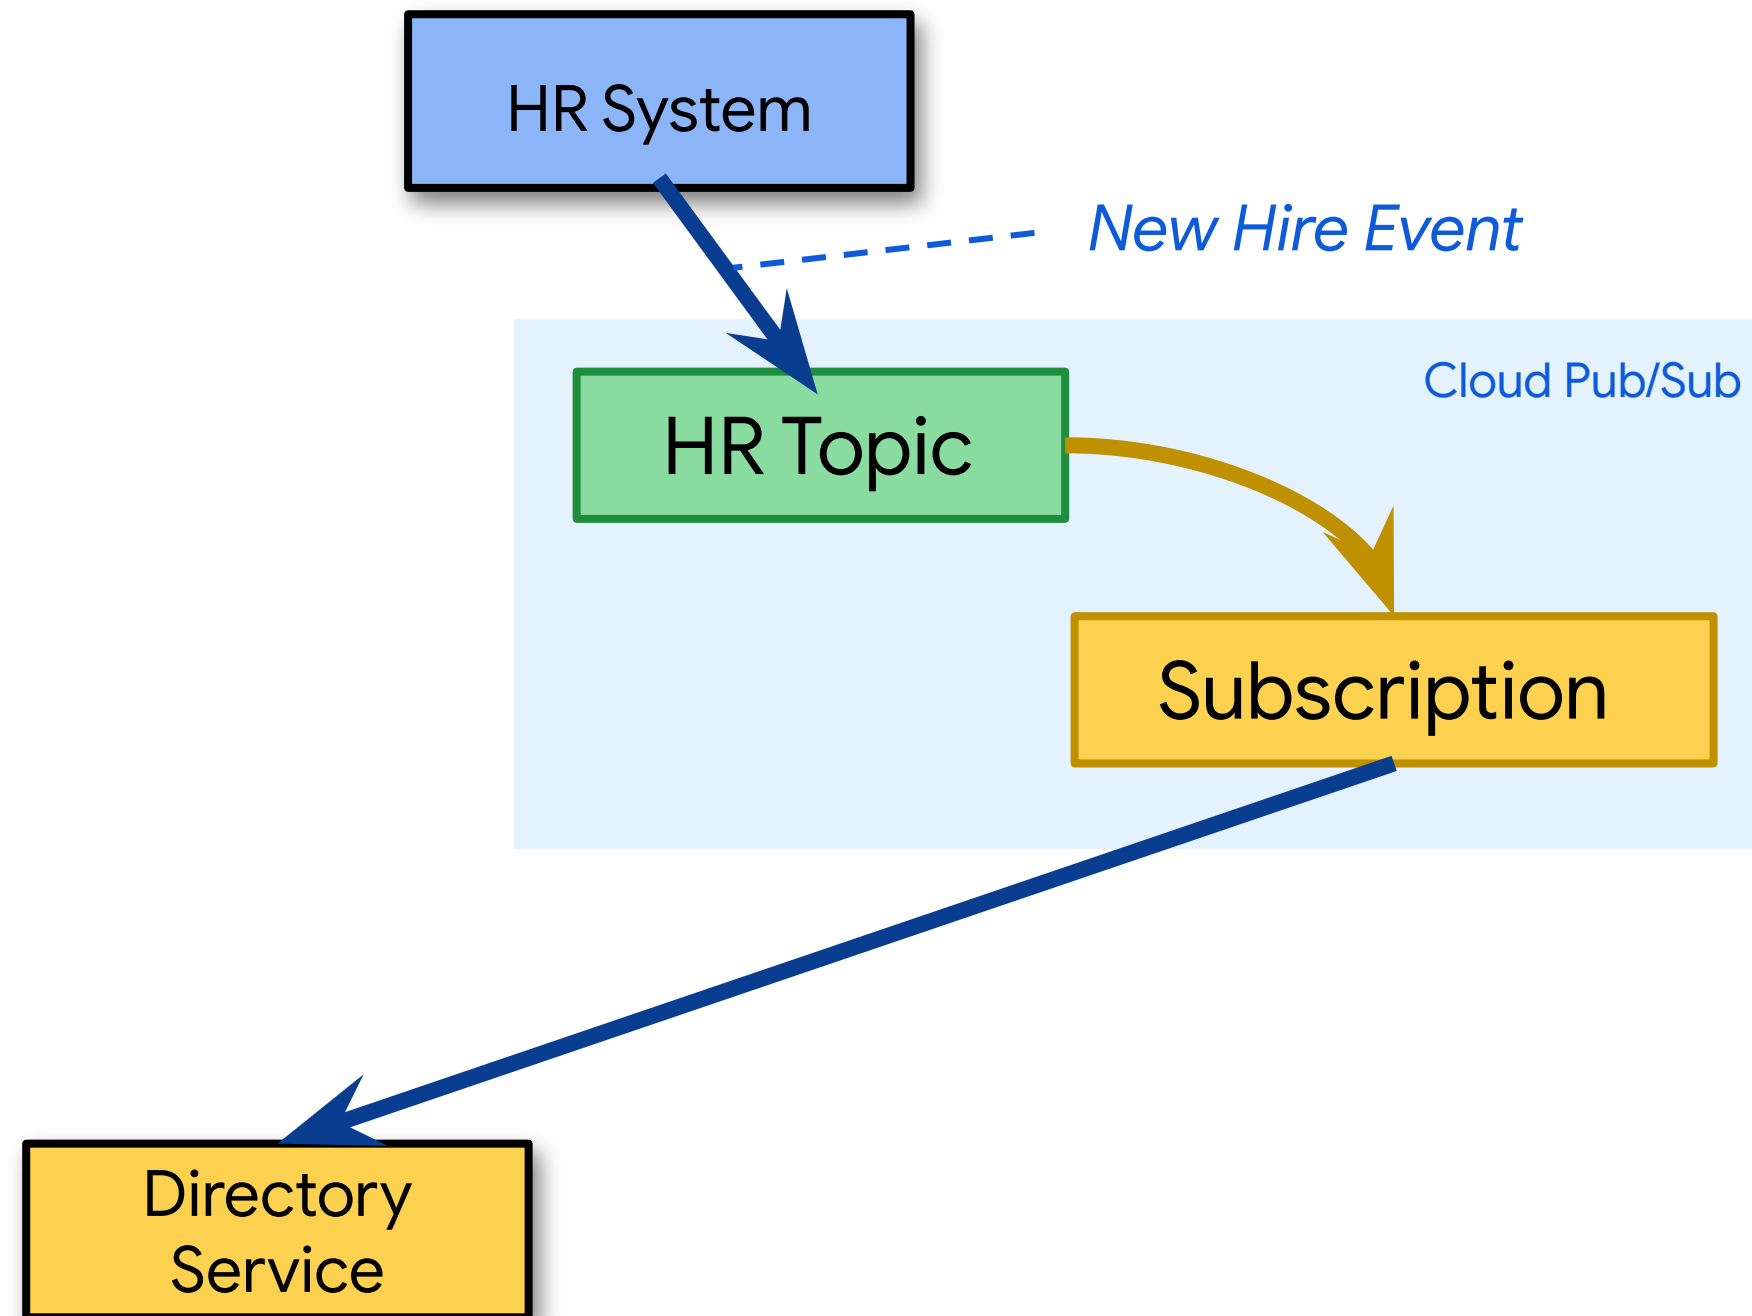
Example of a Cloud Pub/Sub application



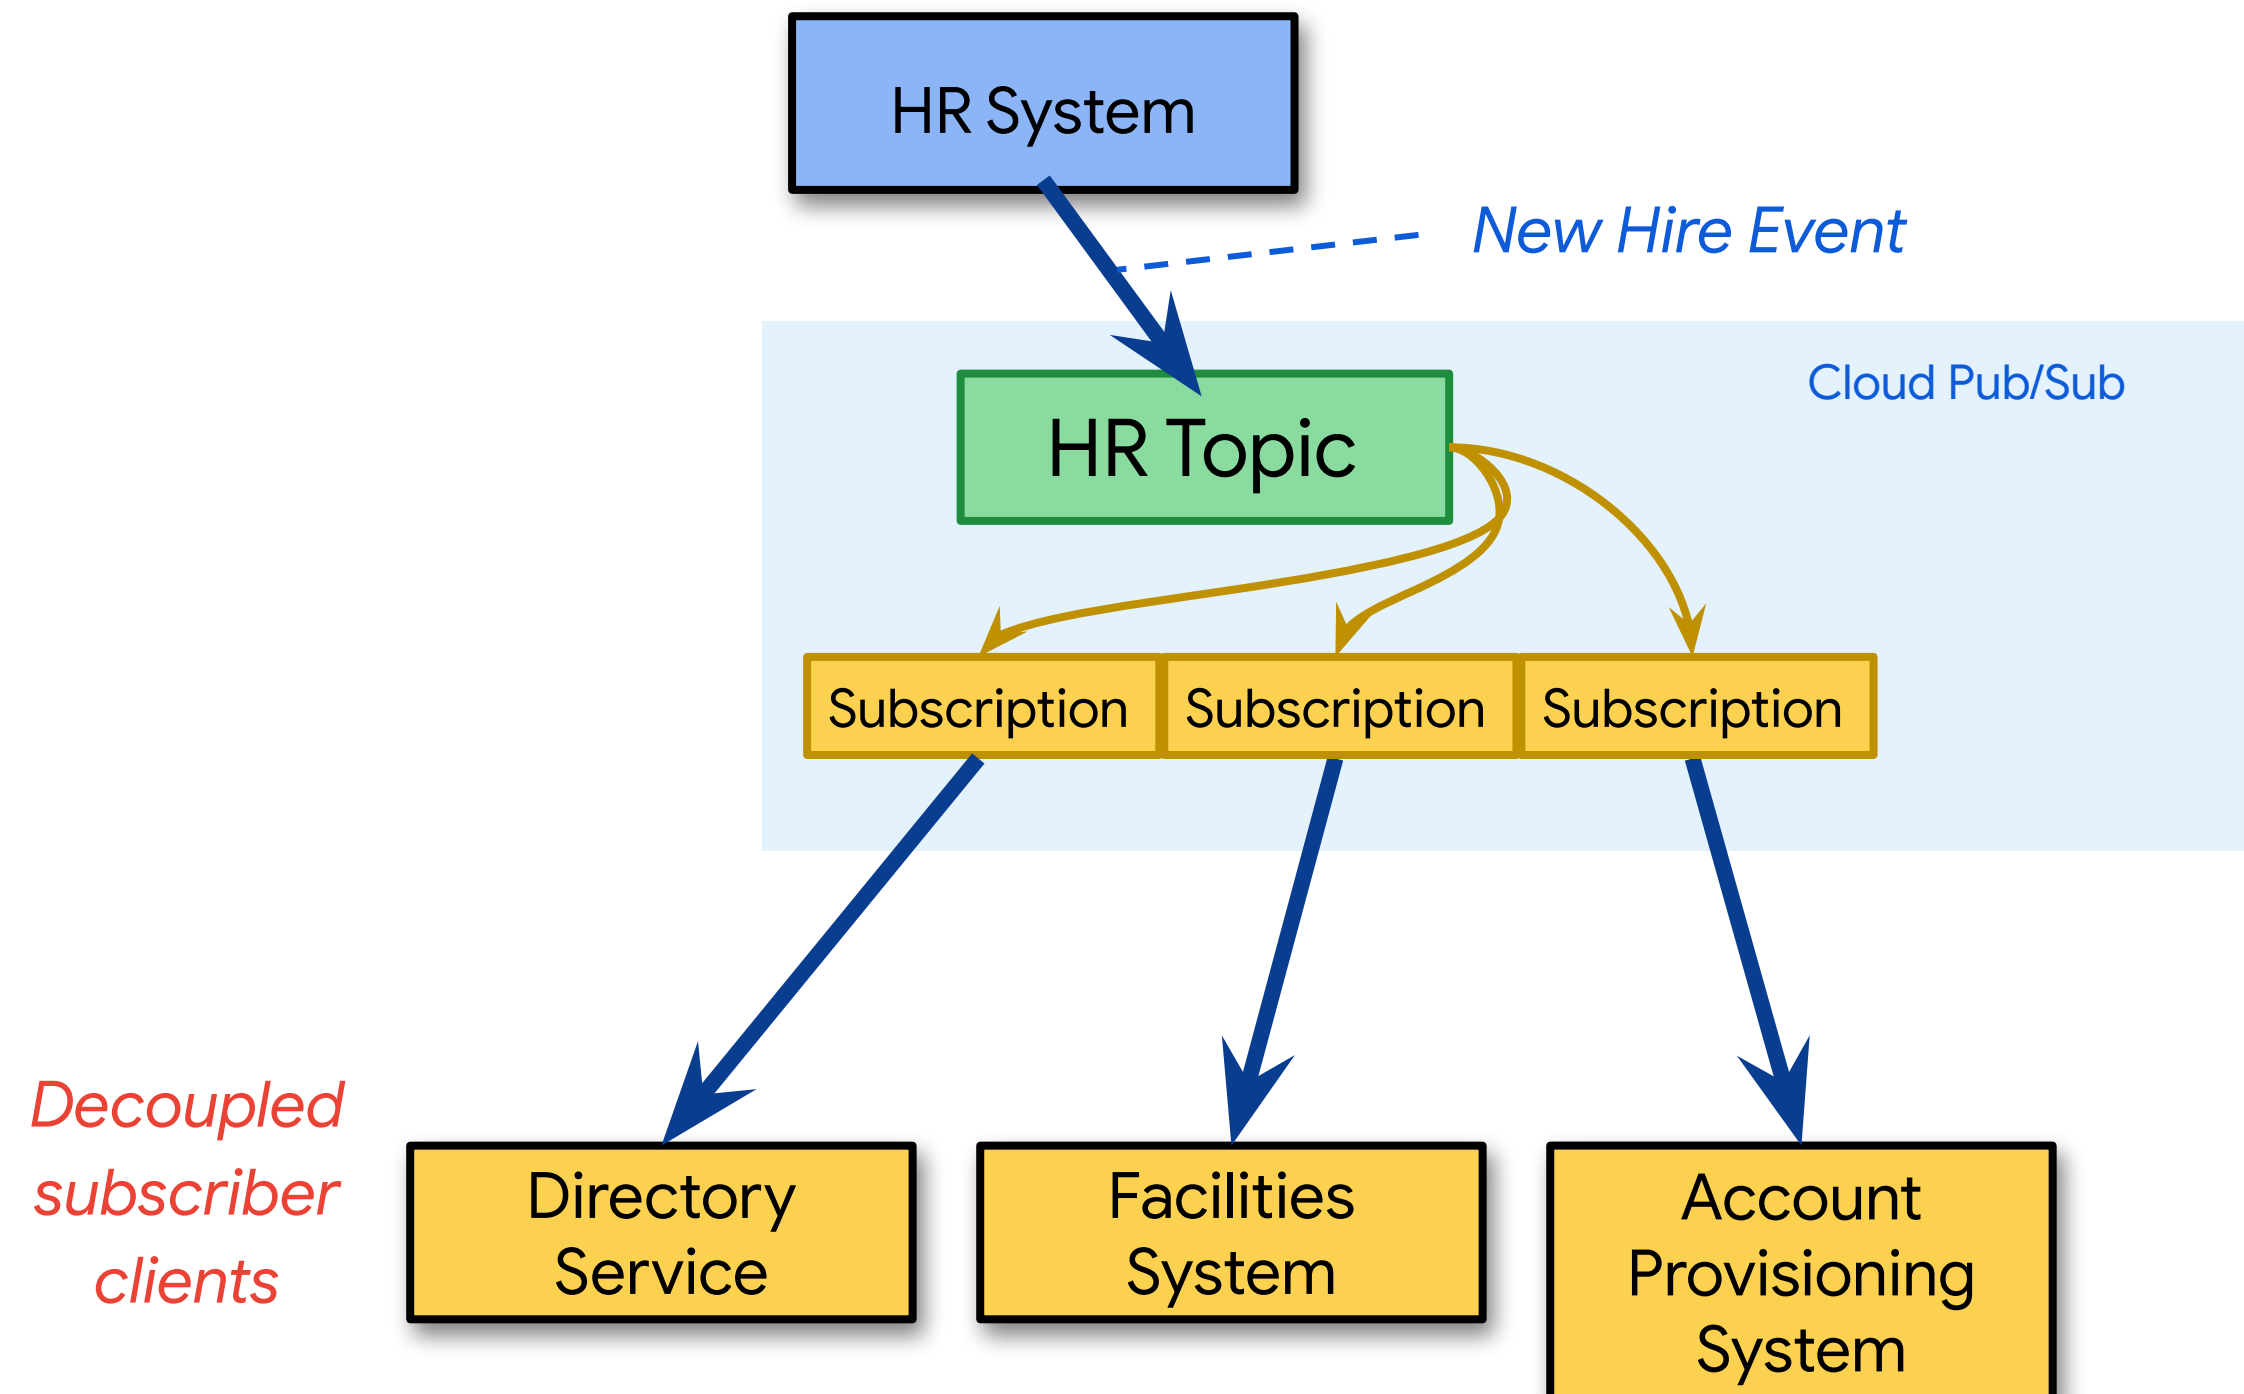
A new employees arrives causing a new hire event.



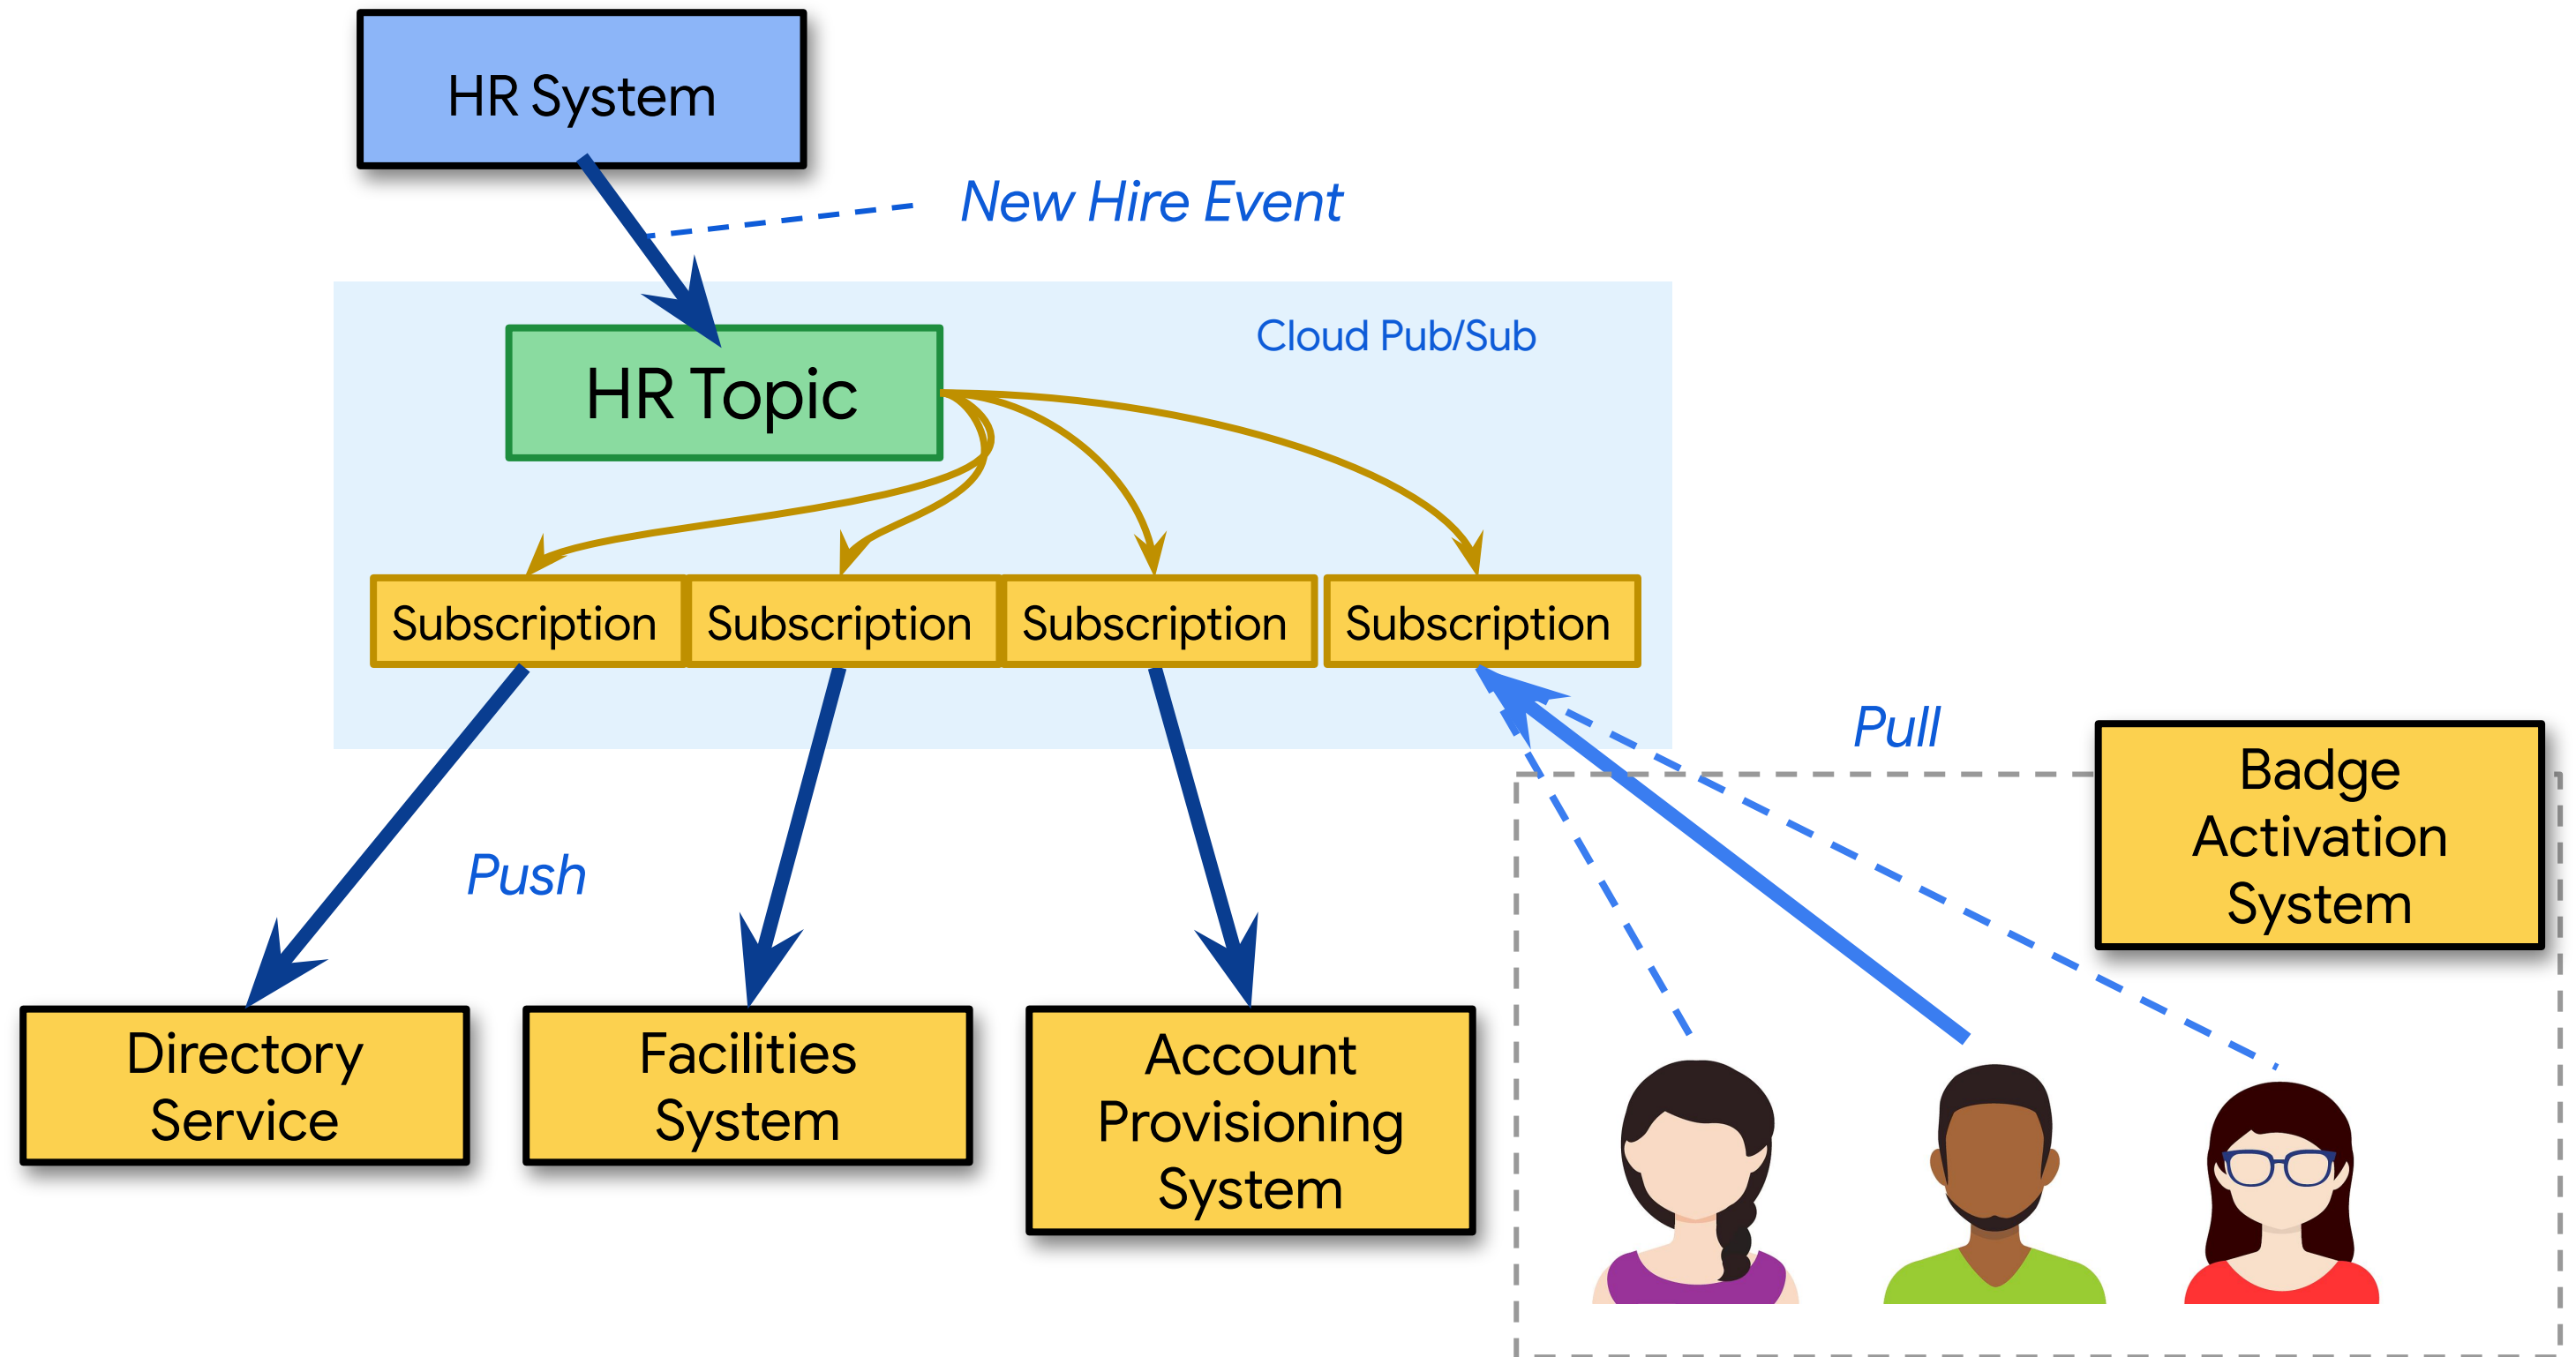
The message is sent from Topic to Subscription



There can be multiple Subscriptions for each Topic

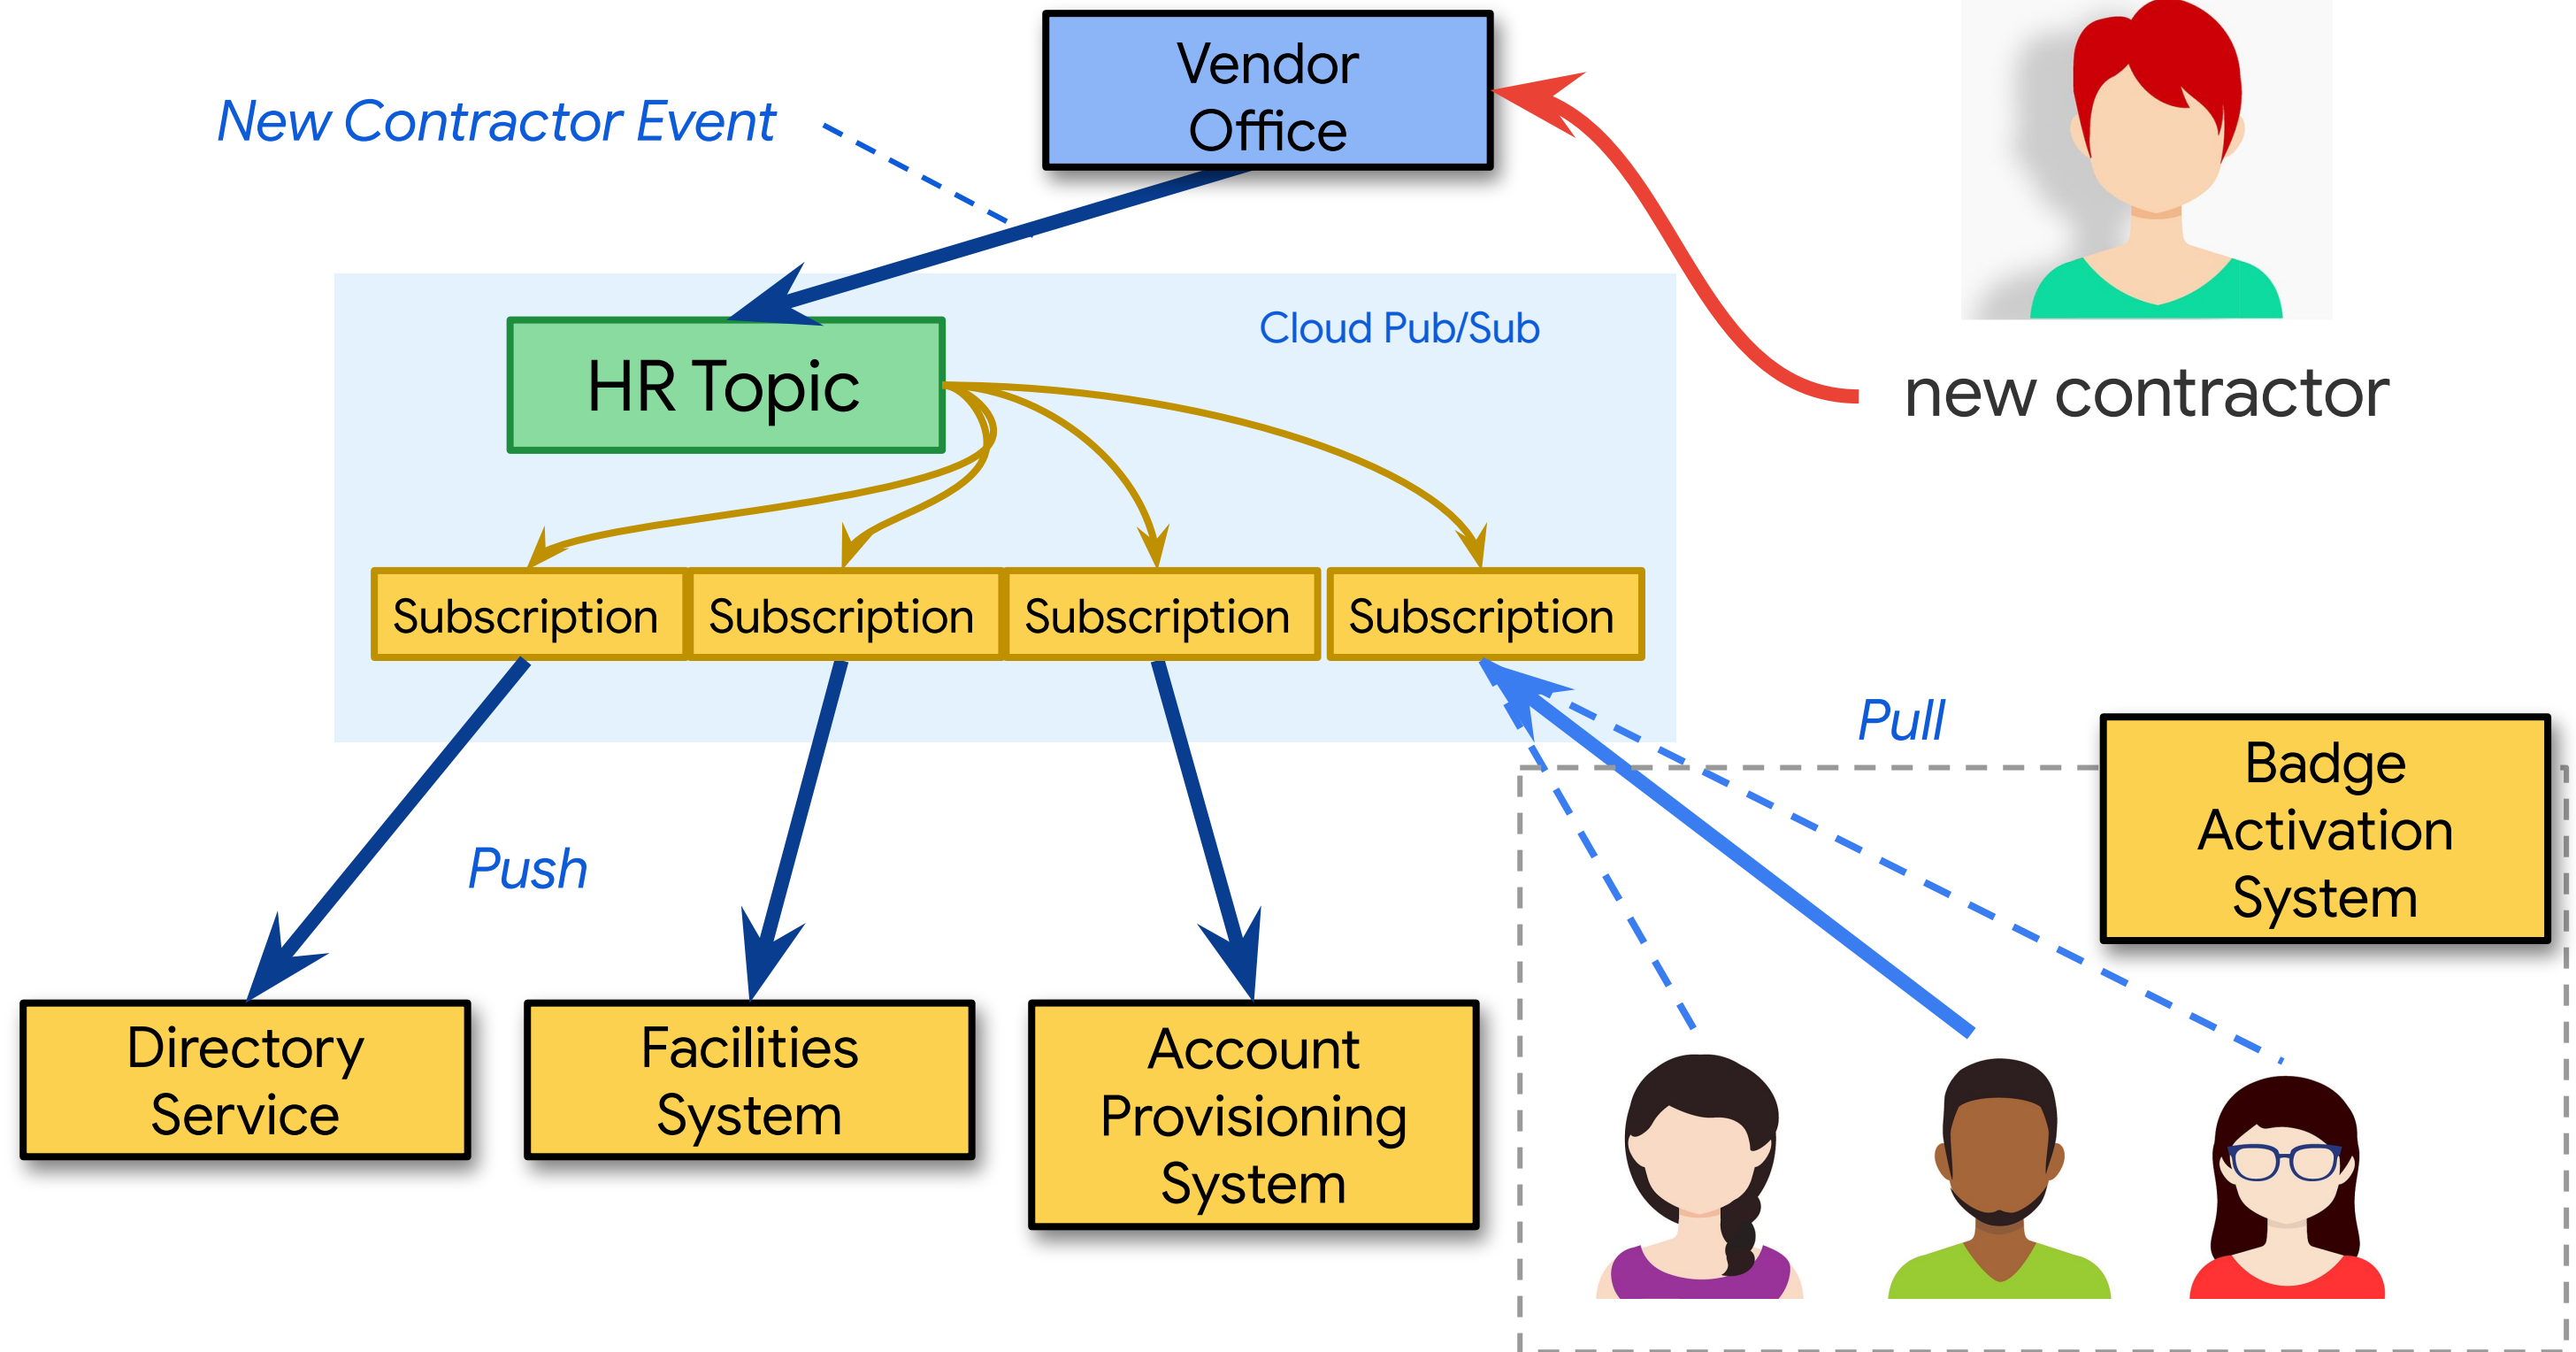


And there can be multiple subscribers per Subscription

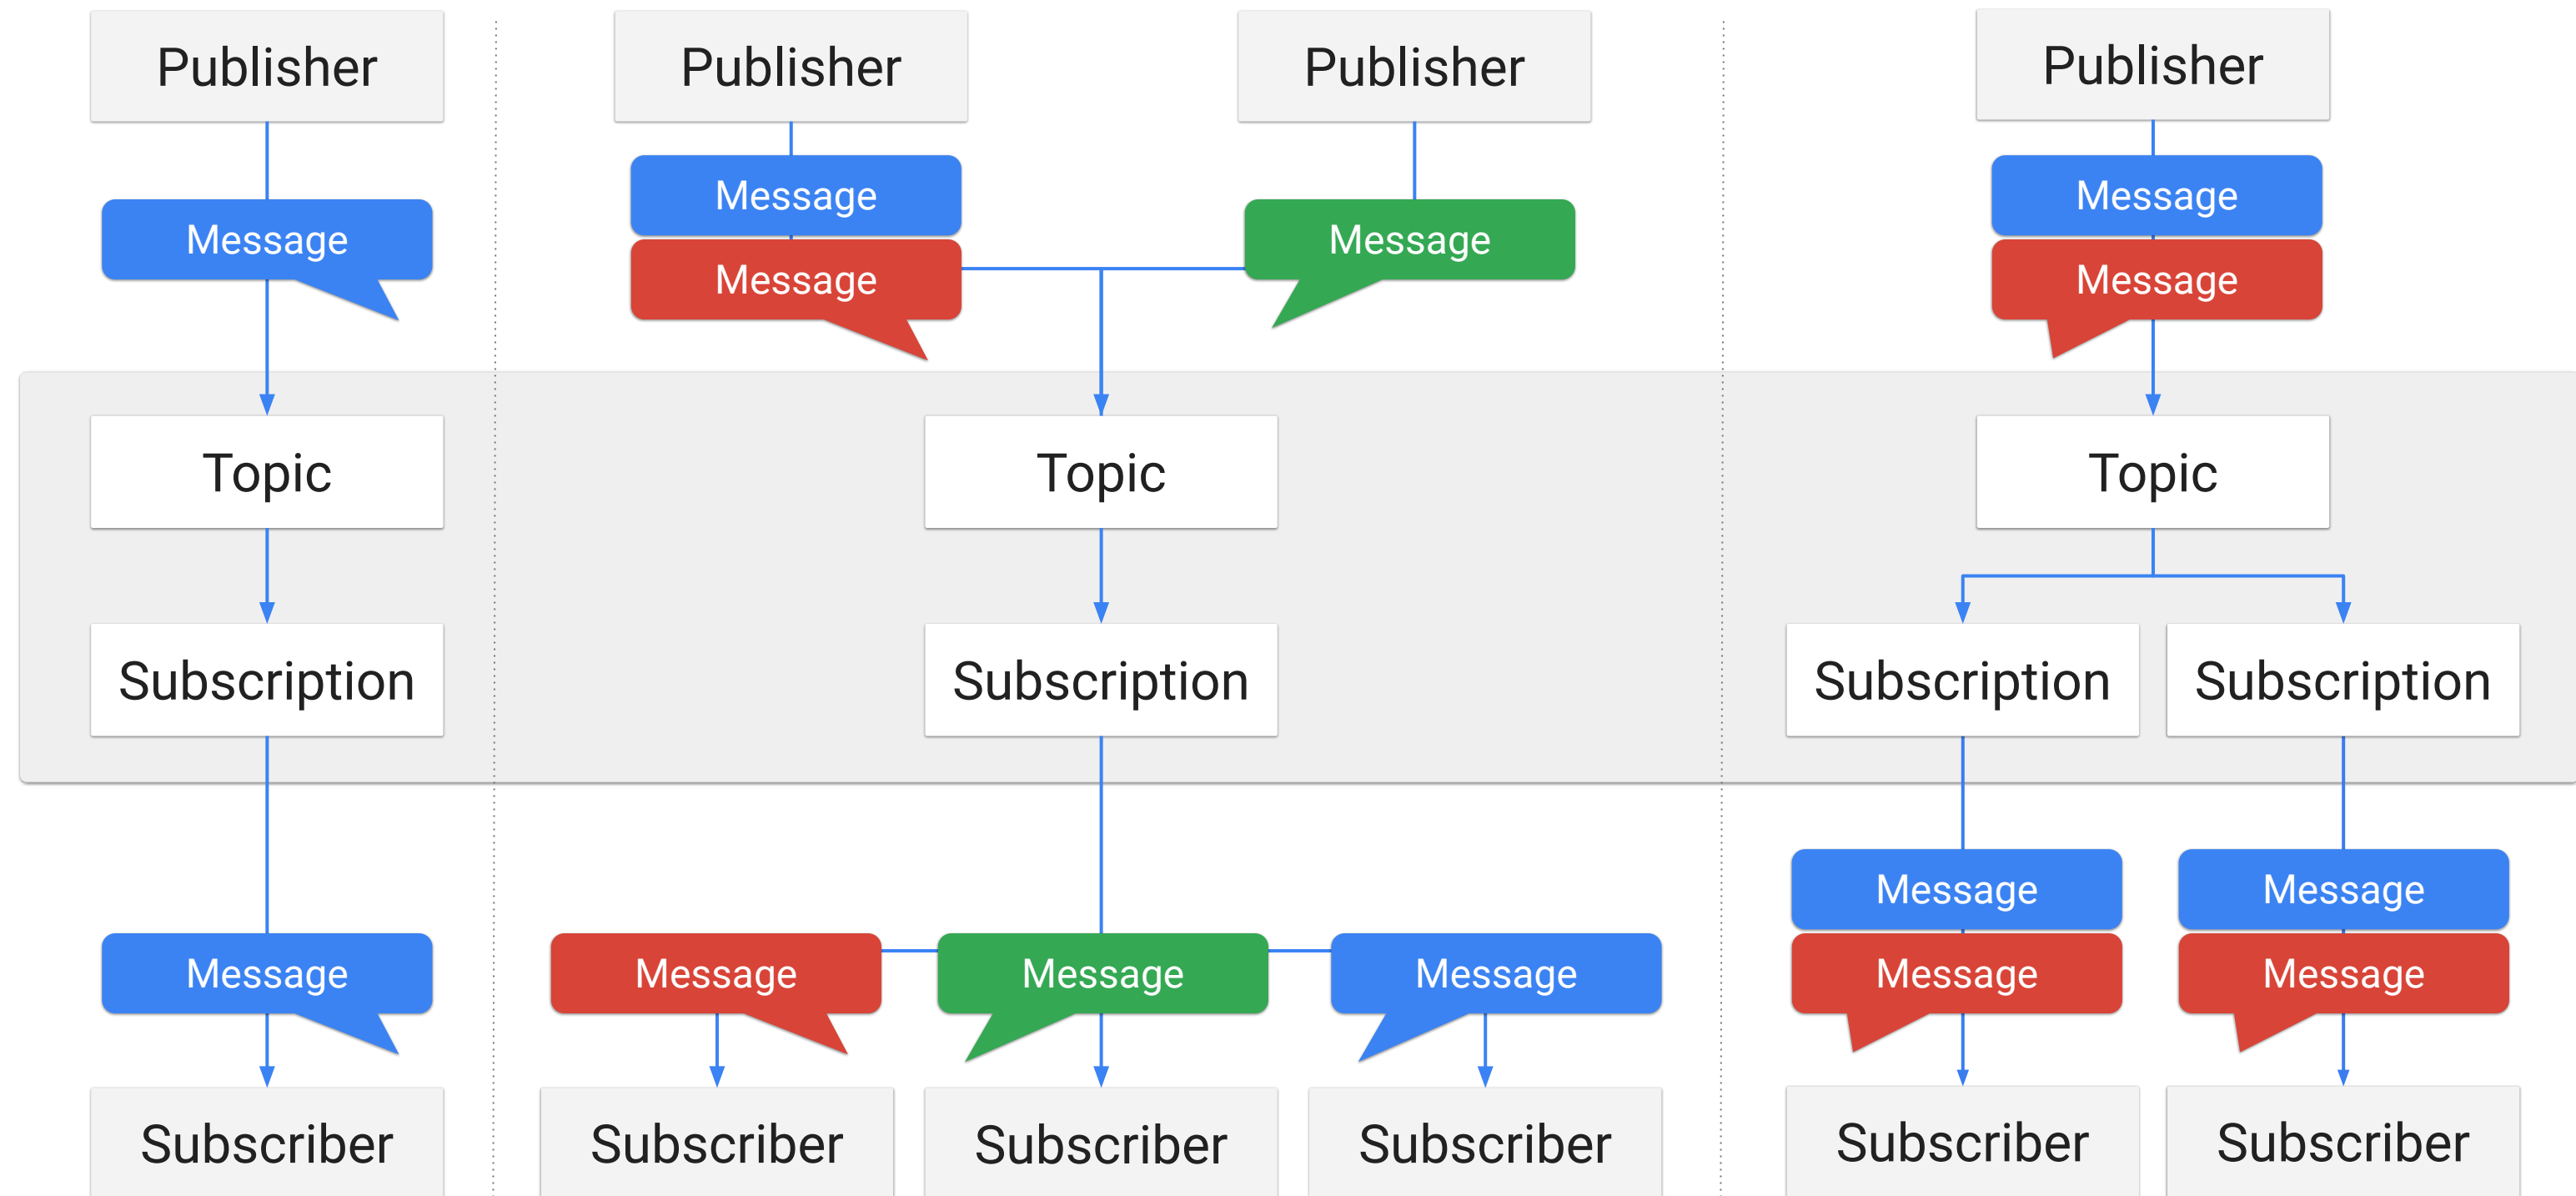


And there can be multiple publishers to the Topic

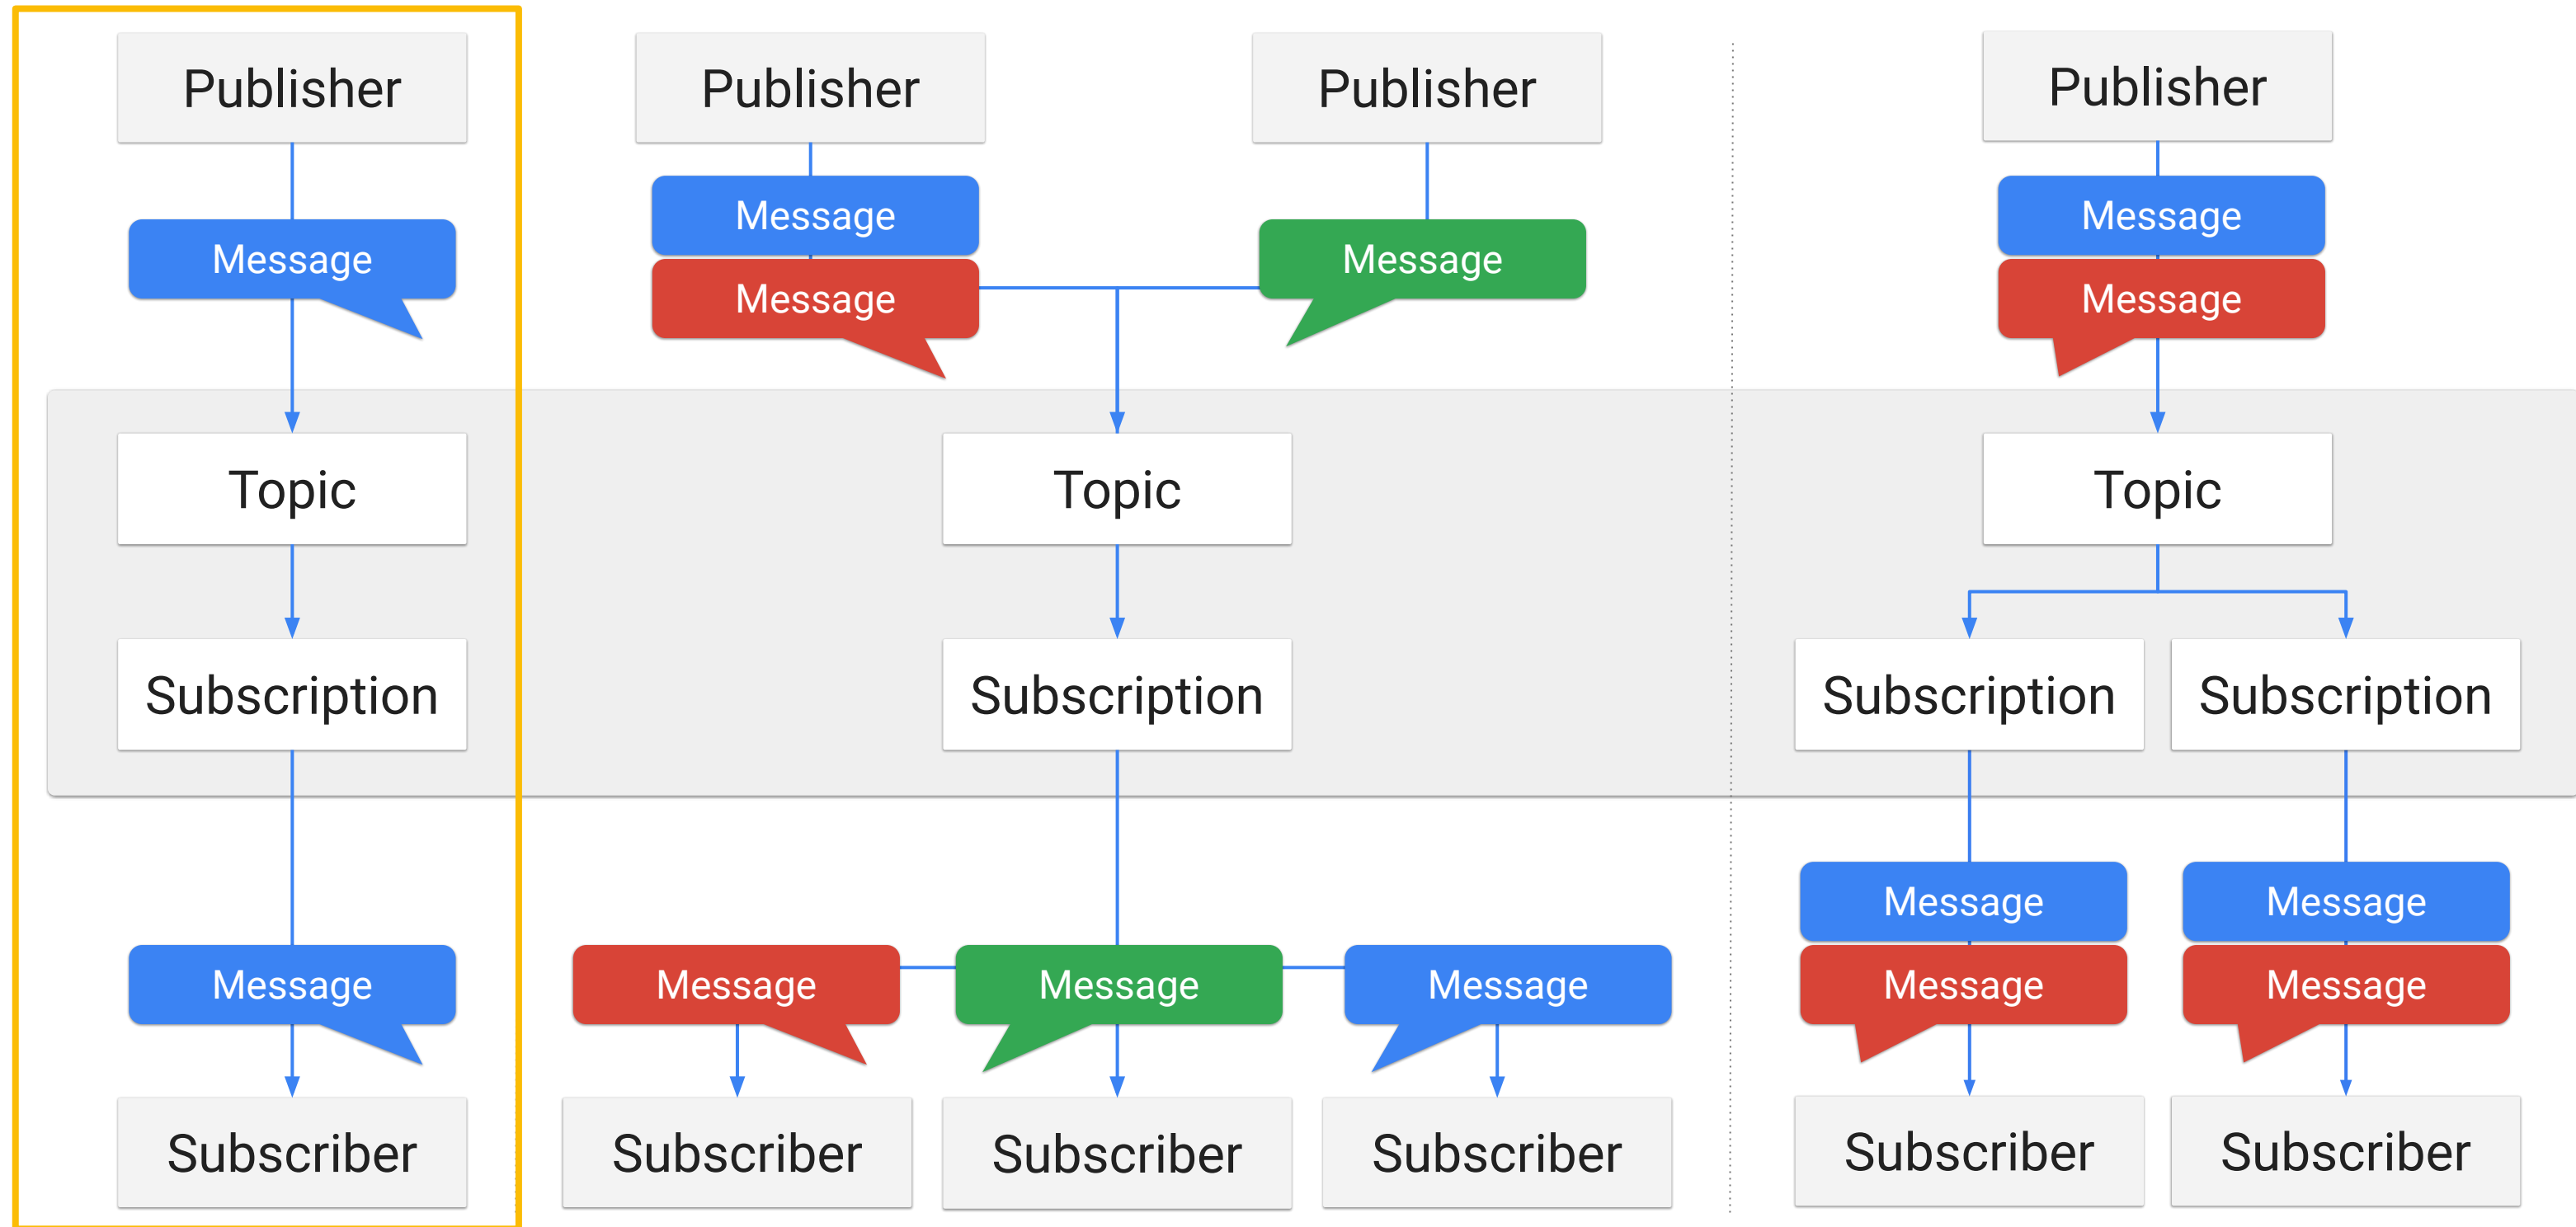
*Decoupled
publisher
clients*



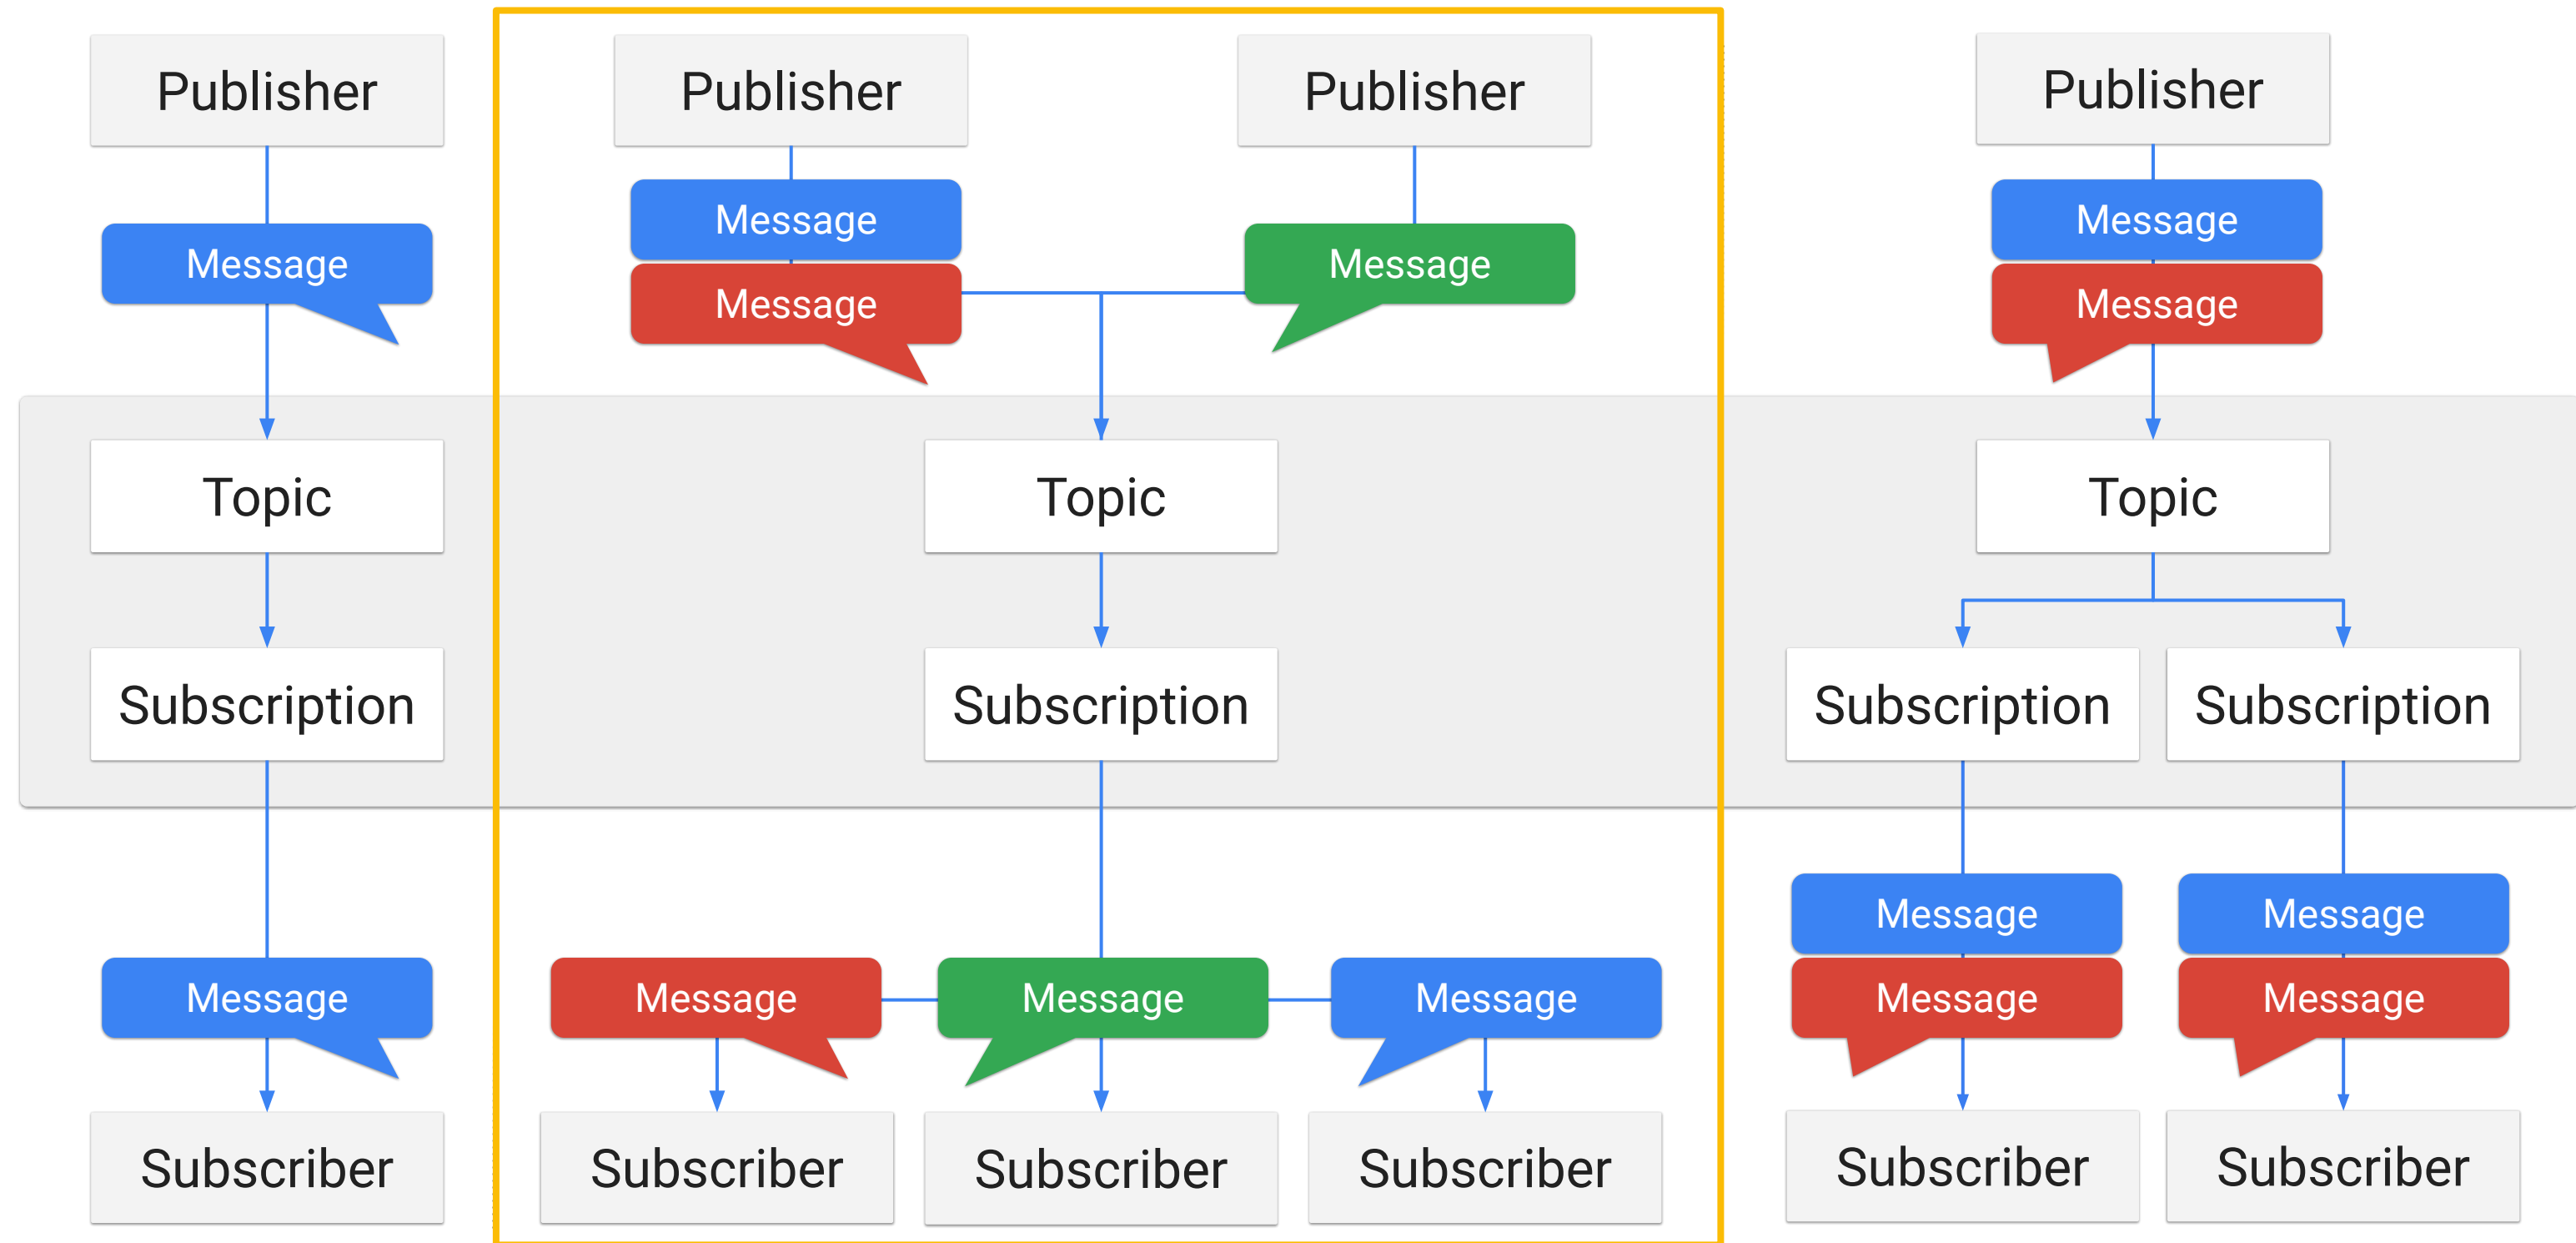
Publish/Subscribe patterns



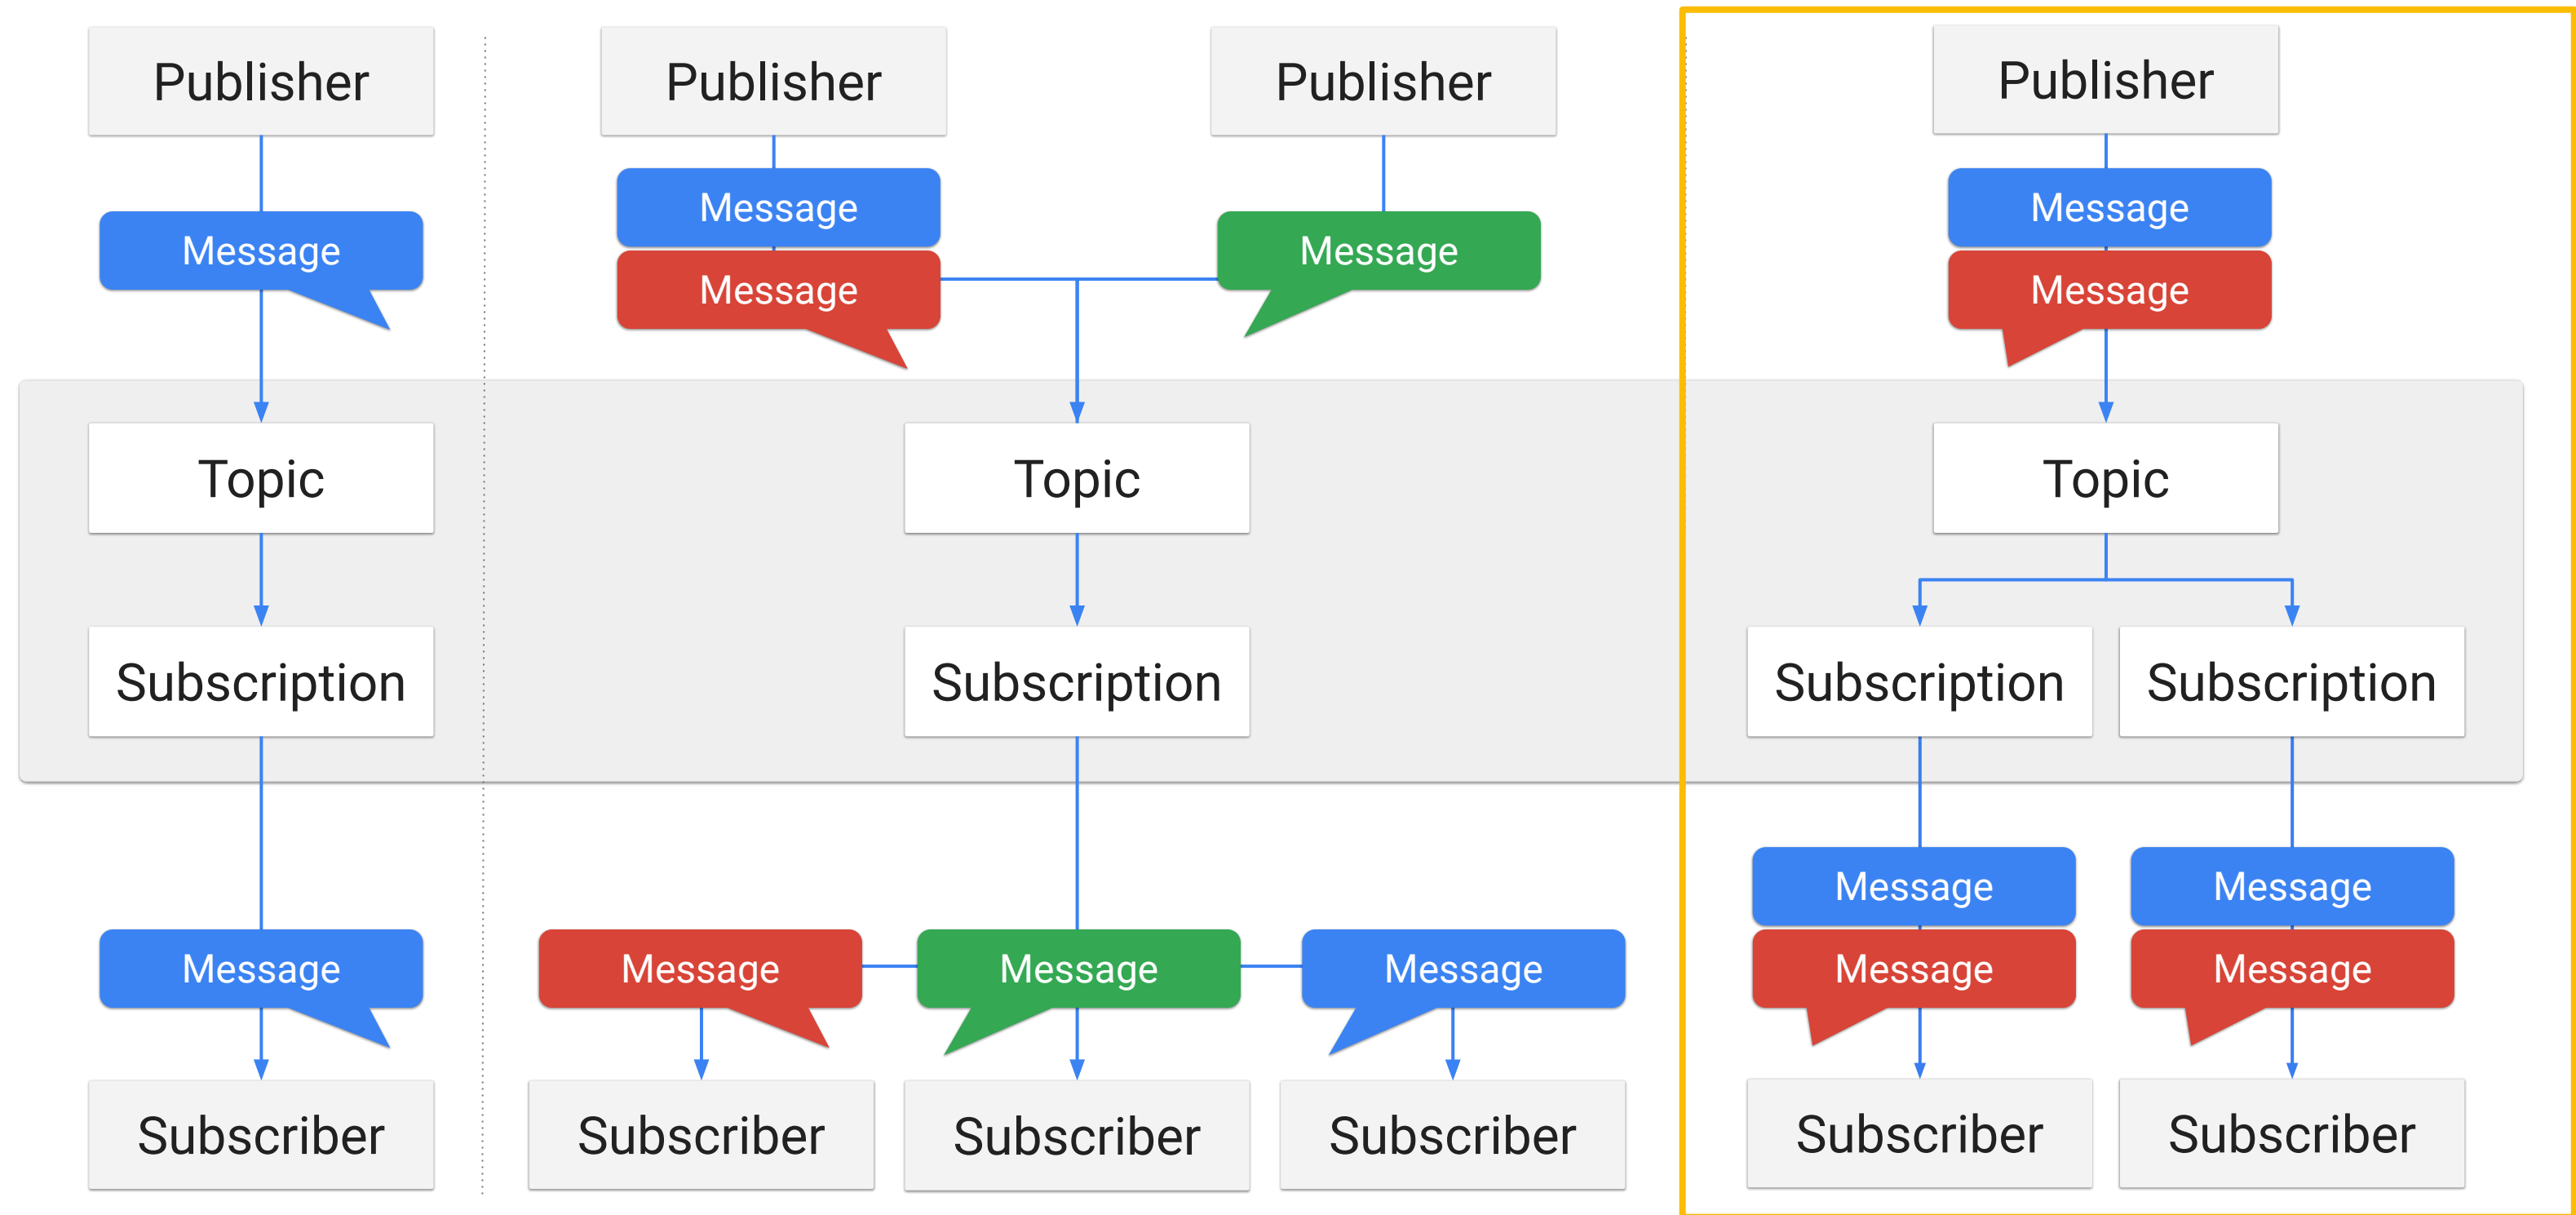
Publish/Subscribe patterns



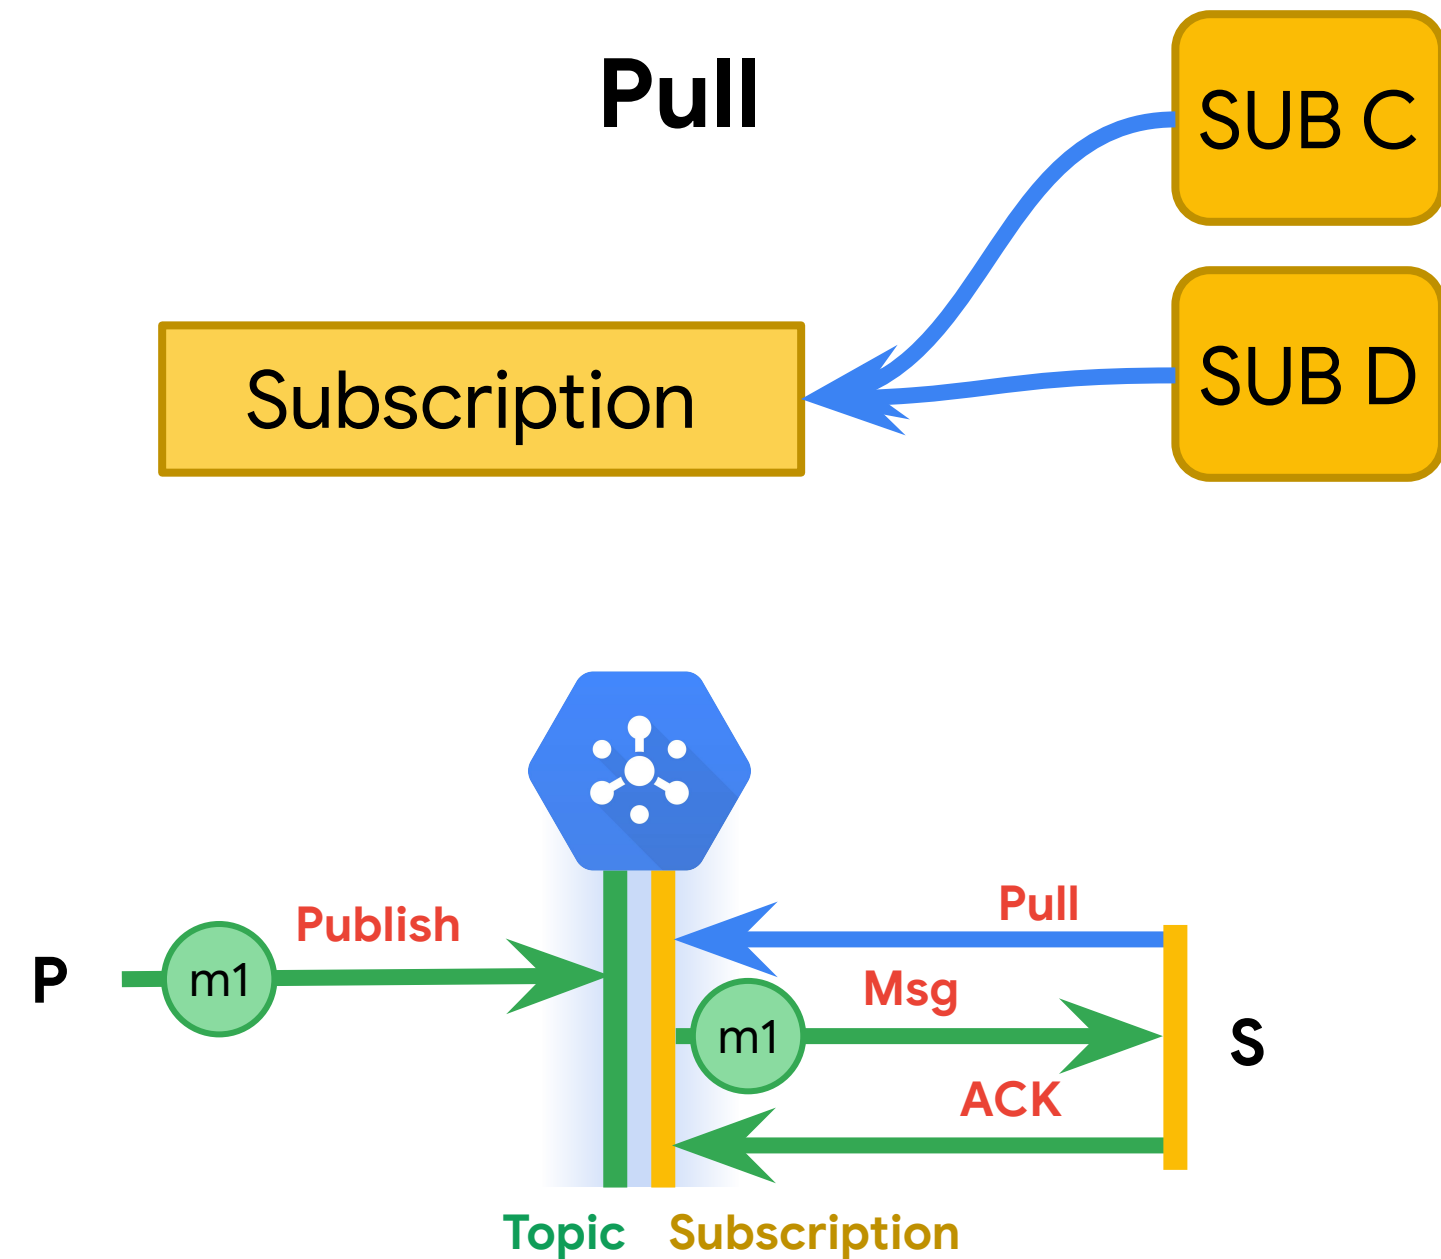
Publish/Subscribe patterns



Publish/Subscribe patterns

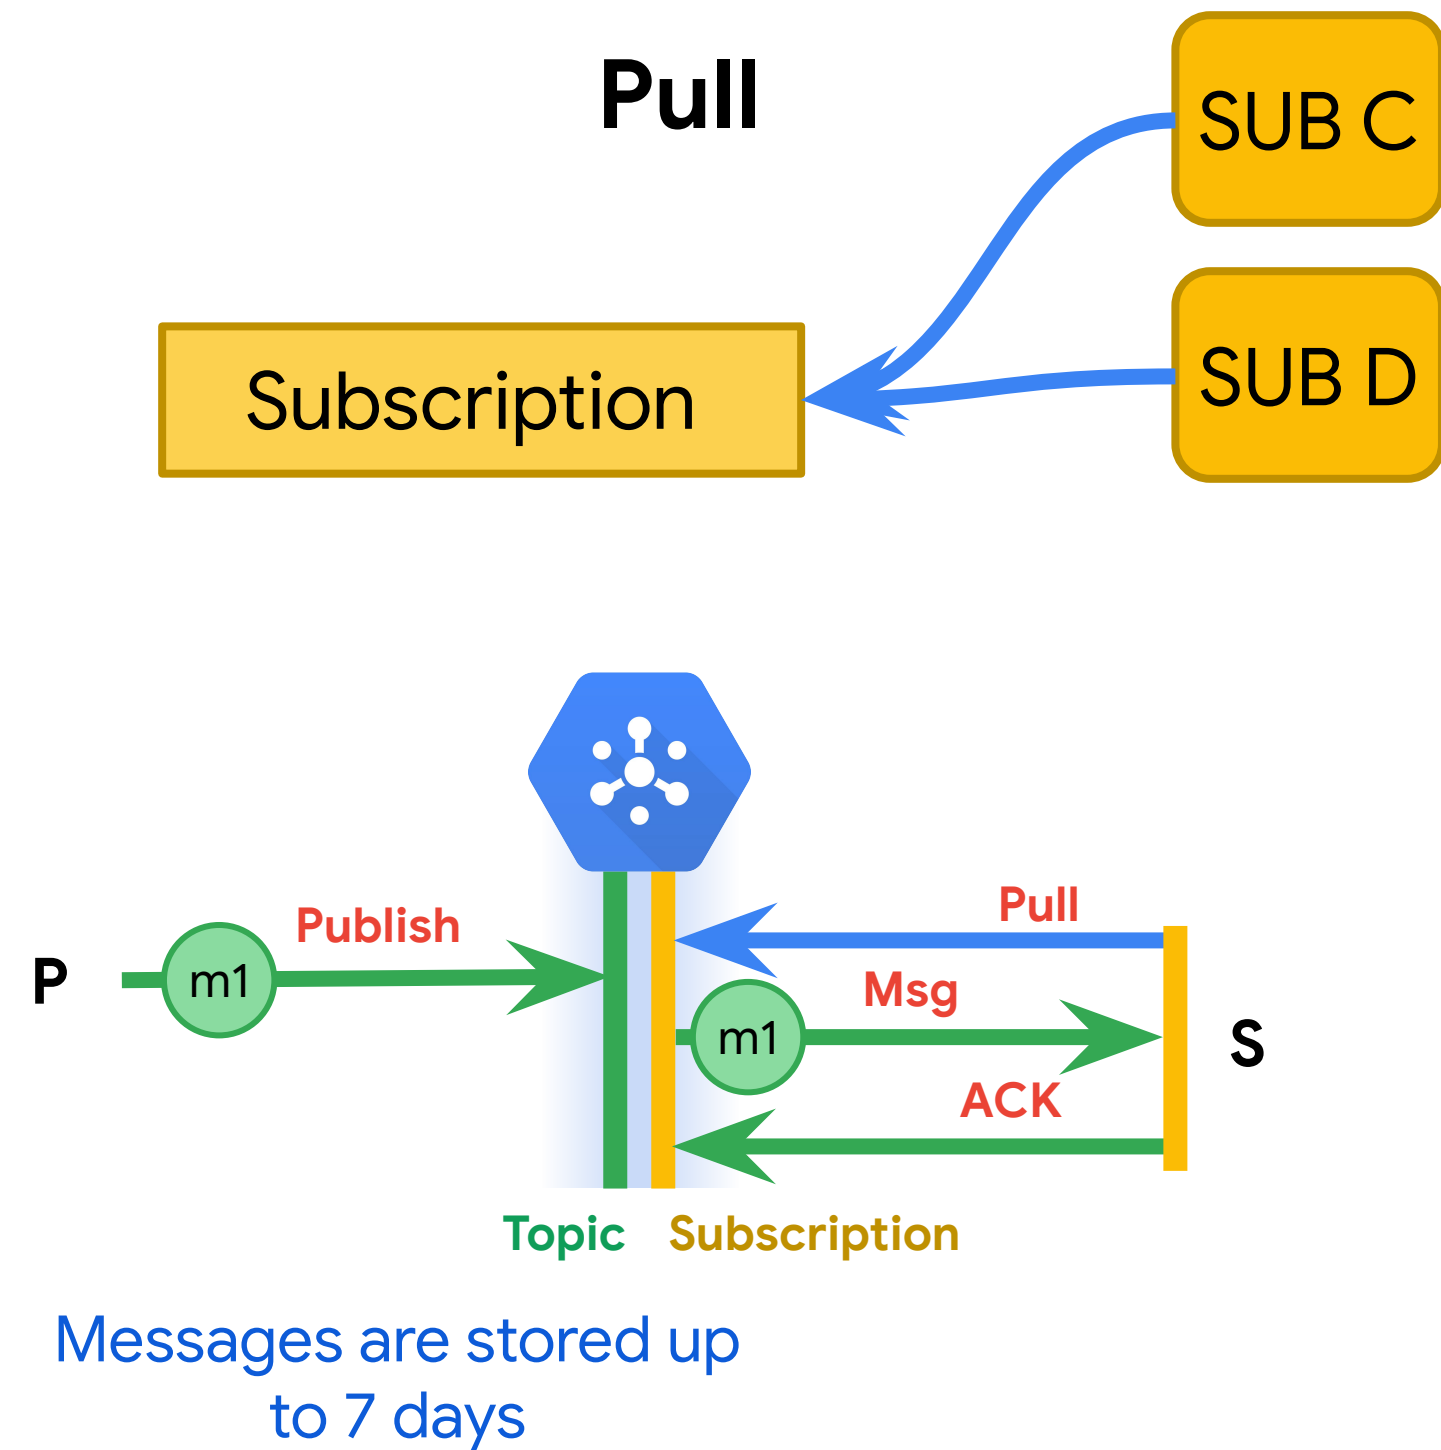
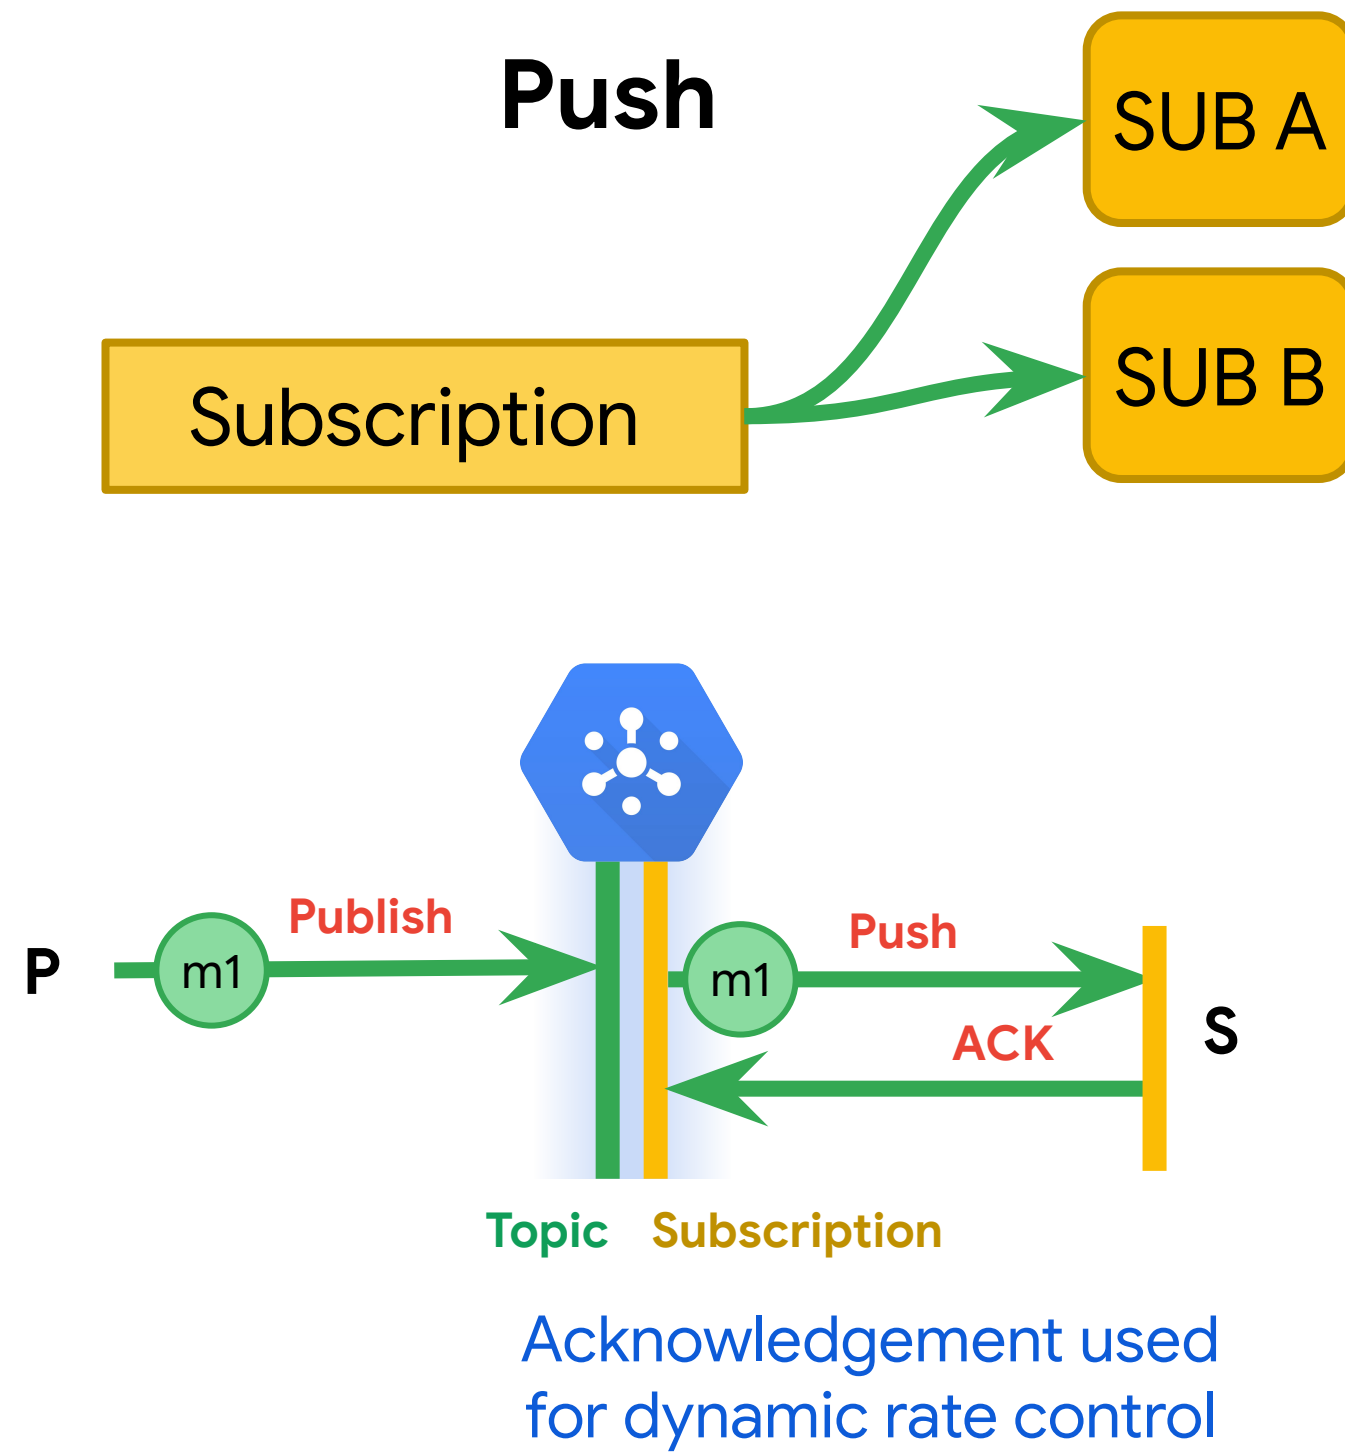


Cloud Pub/Sub provides both Push and Pull delivery



Messages are stored up
to 7 days

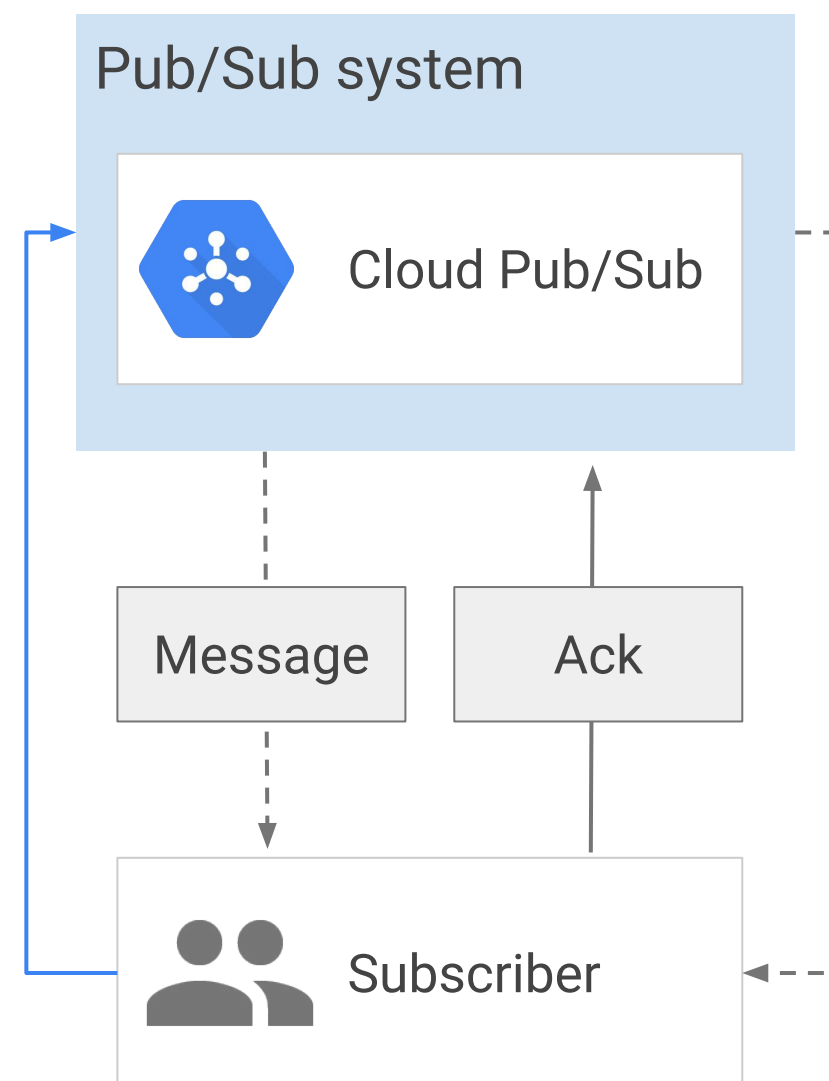
Cloud Pub/Sub provides both Push and Pull delivery



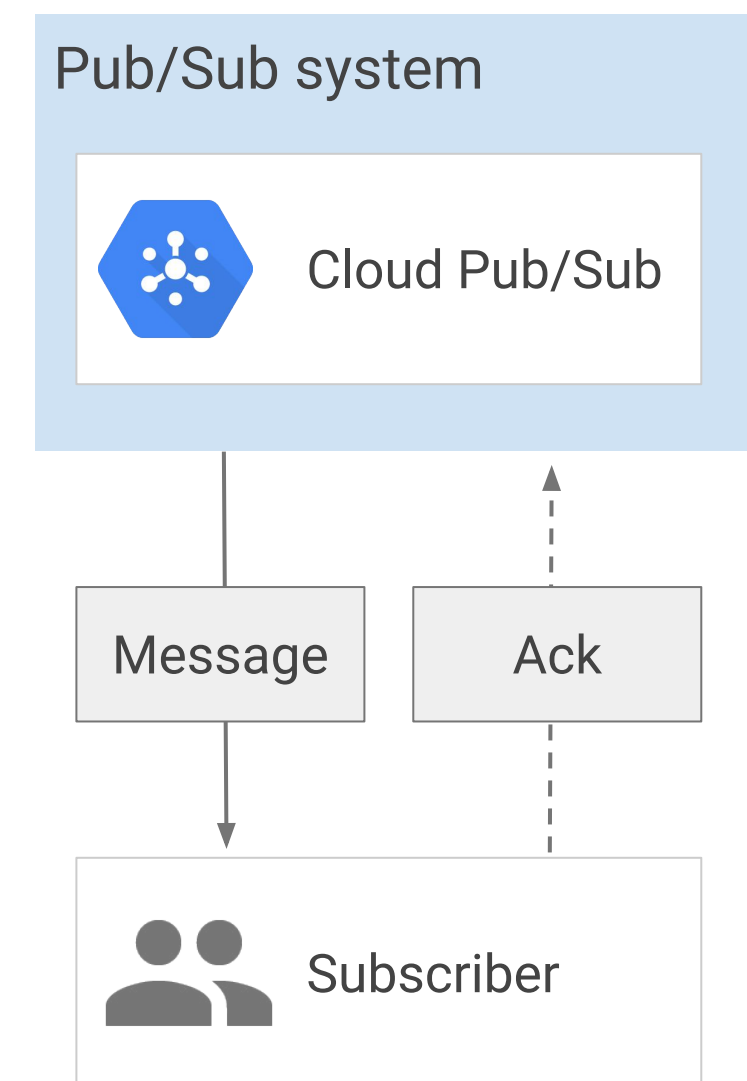
At least once delivery guarantee

- A subscriber ACKs each message for every subscription
- A message is resent if subscriber takes more than **ackDeadline** to respond
- Messages are stored for up to 7 days
- A subscriber can extend the deadline per message

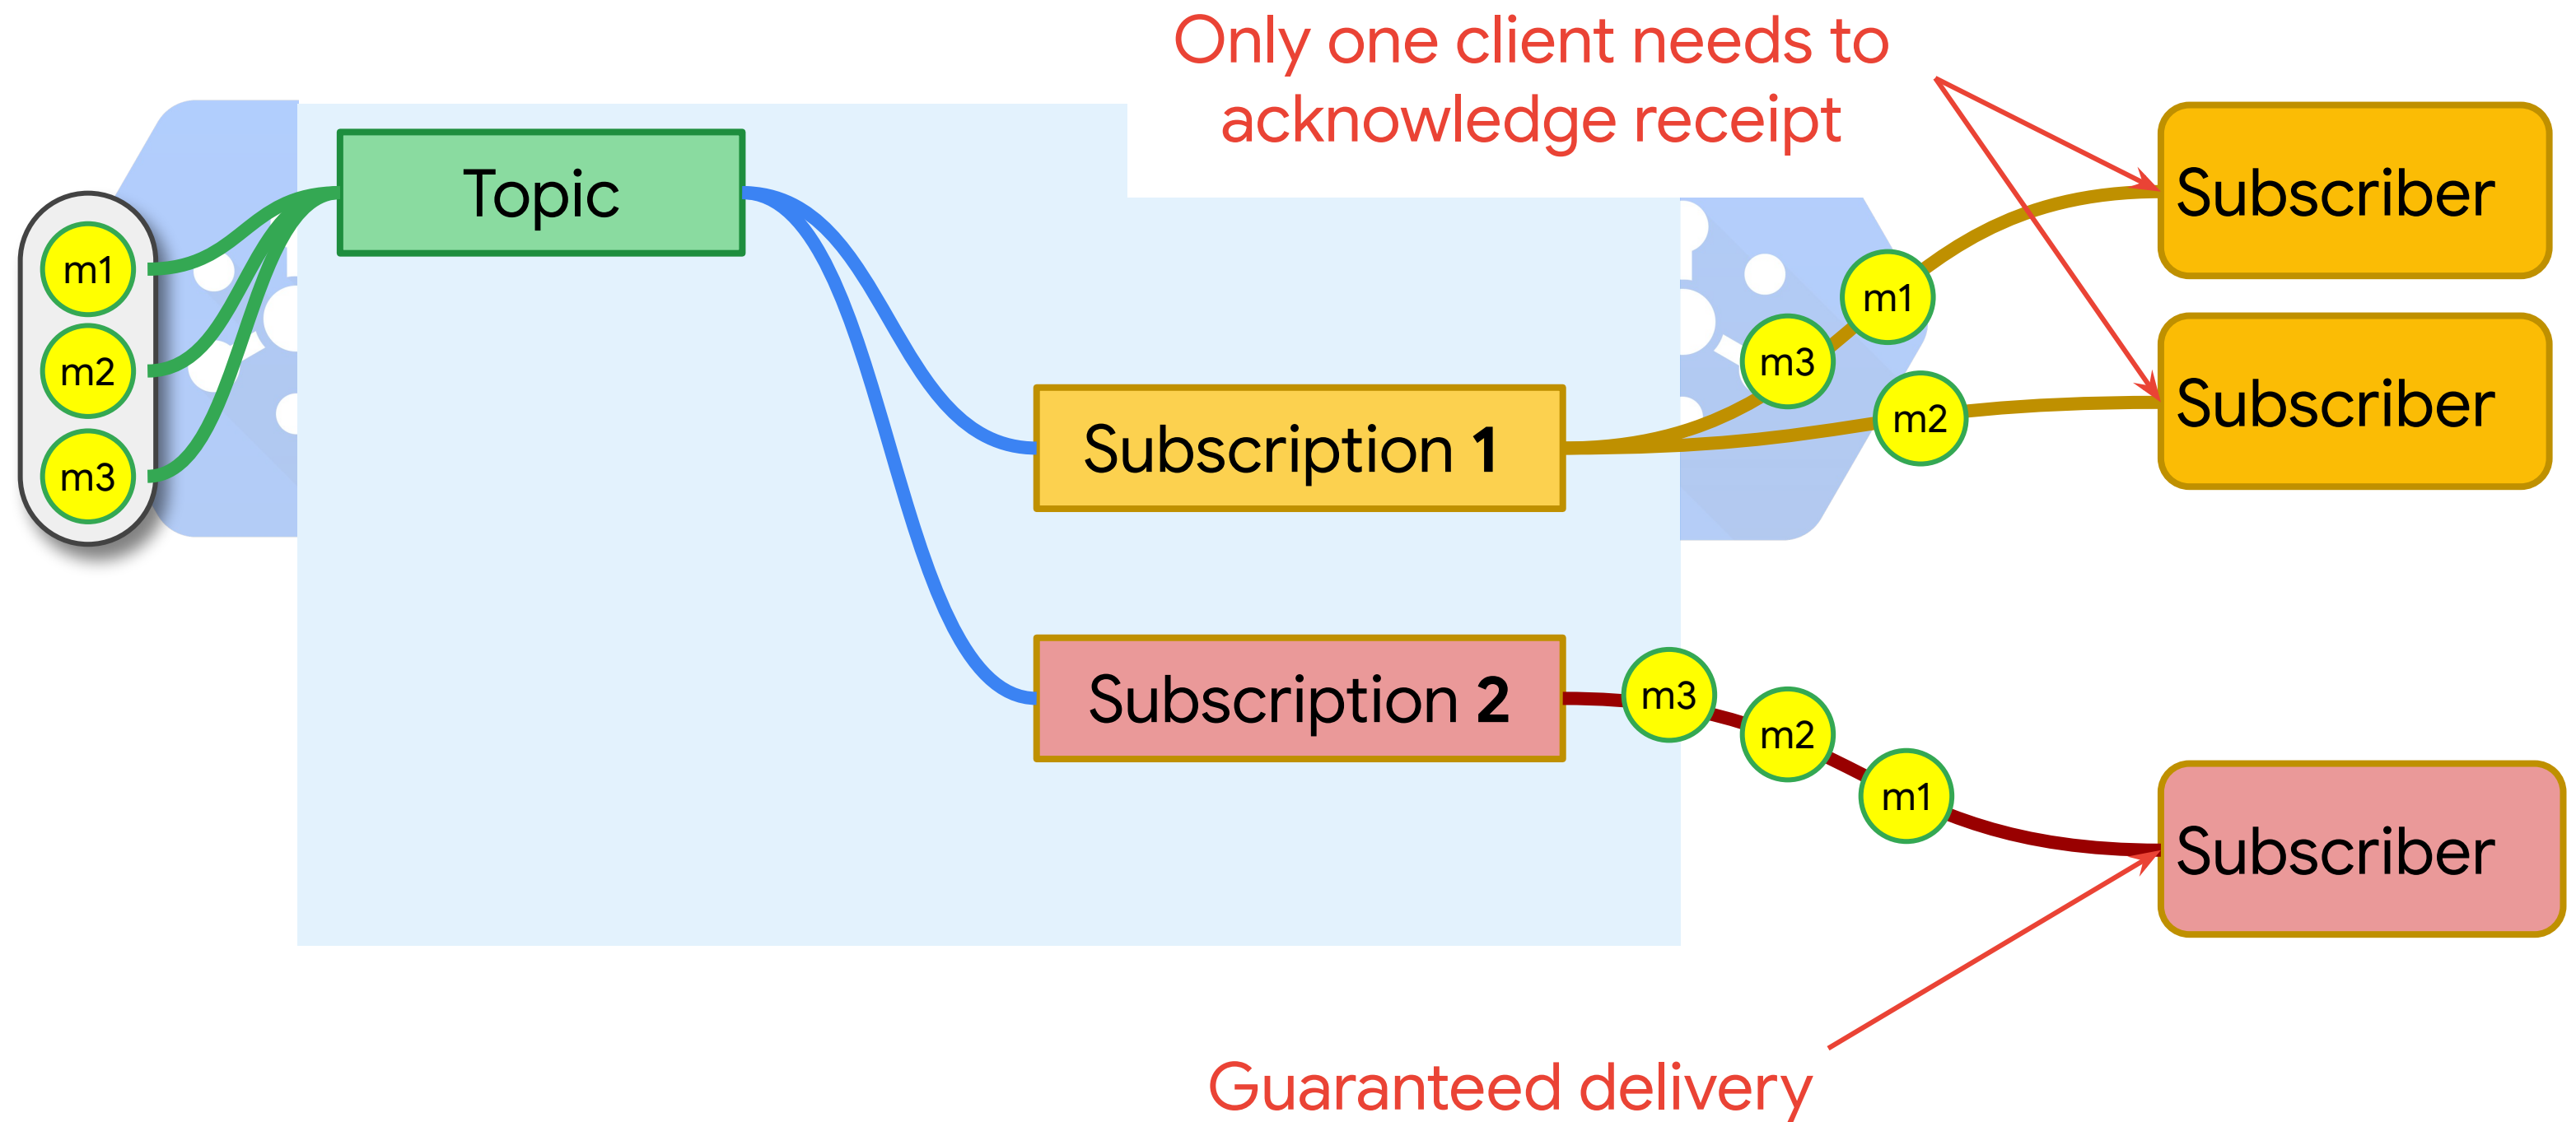
Pull subscription



Push subscription



Subscribers can work as a team or separately



Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

[Create topic](#)

Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego --message  
"hello"
```

Publish to topic

Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego --message  
"hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1

publisher = pubsub_v1.PublisherClient()

topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)

publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

----- Set topic name

Message

Send attribute

Create a client

Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego --message  
"hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1

publisher = pubsub_v1.PublisherClient()

topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)

publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

----- Set topic name

Message

Send attribute

Create a client

Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego --message  
"hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1

publisher = pubsub_v1.PublisherClient()

topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)

publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

----- Set topic name

Message

Send attribute

Create a client

Subscribing with Cloud Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1
```

Python

```
subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME',
)
subscriber.create_subscription(
    name=subscription_name, topic=topic_name)

def callback(message):
    print(message.data)
    message.ack()

future = subscriber.subscribe(subscription_name, callback)
```

Select
topic
name

Set subscription
name

callback when
message received

Create a client

Push method
Callback function

Subscribing with Cloud Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1
```

Python

```
subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME',
)
subscriber.create_subscription(
    name=subscription_name, topic=topic_name)
```

```
def callback(message):
    print(message.data)
    message.ack()
```

```
future = subscriber.subscribe(subscription_name, callback)
```

Create a client

Select
topic
name

Set subscription
name

callback when
message received

Push method
Callback function

Subscribing with Cloud Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1
```

Python

```
subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME',
)
subscriber.create_subscription(
    name=subscription_name, topic=topic_name)
```

Select
topic
name

Set subscription
name

```
def callback(message):
    print(message.data)
    message.ack()

future = subscriber.subscribe(subscription_name, callback)
```

callback when
message received

Create a client

Push method
Callback function

Subscribing with Cloud Pub/Sub using synchronous pull

```
gcloud pubsub subscriptions create --topic sandiego mySub1
```

Create subscription

```
gcloud pubsub subscriptions pull --auto-ack mySub1
```

Pull subscription

```
import time

from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()

subscription_path = subscriber.subscription_path(project_id, subscription_name)

NUM_MESSAGES = 2
ACK_DEADLINE = 30
SLEEP_TIME = 10

# The subscriber pulls a specific number of messages.
response = subscriber.pull(subscription_path, max_messages=NUM_MESSAGES)
```

Set subscription name

Create a client

``projects/{project_id}/subscriptions/{subscription_name}`` ← subscription_path format

Subscriber is non-blocking

Keep the main thread from exiting to allow it to process messages synchronously

By default, the Publisher batches messages; turn this off if you desire lower latency



Batching messages: throughput versus latency

Changing the batch settings in Cloud Pub/Sub

Python

```
from google.cloud import pubsub
from google.cloud.pubsub import types

client = pubsub.PublisherClient(
    batch_settings=BatchSettings(max_messages=500),
)
```


[Change batch setting](#)

Pub/Sub: latency, out-of-order, duplication will happen


- Latency -- no guarantees
- Messages can be delivered in any order, especially with large backlog
- Duplication may happen

Cloud Pub/sub with Dataflow: Exactly once, ordered processing



- ✓ Cloud Pub/Sub delivers at least once
- ✓ Cloud Dataflow: Deduplicate, order, and window
- ✓ Separation of concerns → scale



Use Cloud Pub/Sub for streaming resilience





Use Cloud Pub/Sub for streaming resilience



Security, monitoring, and logging for Pub/Sub

Cloud Logging metrics

pubsub_topic
pubsub_subscription



Cloud Audit Logs

Admin = on
Data Access = off

Authentication

Service accounts
User accounts



Access control

Pub/Sub roles per topic



Publish Streaming Data
into Pub/Sub

Lab Objectives

Create a Pub/Sub topic and subscription

Simulate your traffic sensor data into Pub/Sub